



<https://hao-ai-lab.github.io/dsc204a-f25/>

# DSC 204A: Scalable Data Systems

## Fall 2025

---

Staff

Instructor: Hao Zhang

TAs: Mingjia Huo, Yuxuan Zhang



[@haozhangml](https://twitter.com/haozhangml)



[@haoailab](https://twitter.com/haoailab)



[haozhang@ucsd.edu](mailto:haozhang@ucsd.edu)

# Where We Are

Machine Learning Systems

Big Data

Cloud

Foundations of Data Systems

1980 - 2000

# Logistics

- Beginning of Quarter Survey: 77% completion
- Finish the 3% and you all get 1 point!

# Qualitative Estimates of Locality

Assuming row-major  
array

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];

    return sum;
}
```

Answer: yes

a		a	a		a		a		a
[0]	...	[0]	[1]	...	[1]	...	[M-1]	...	[M-1]
[0]		[N-1]	[0]		[N-1]		[0]		[N-1]

**Question:** Does this function have good locality with respect to array a?

# Locality Example

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```

**Answer: no, unless...**

**M is very small**

- **Question:** Does this function have good locality with respect to array *a*?

a		a	a		a		a		a
[0]	...	[0]	[1]	...	[1]	...	[M-1]	...	[M-1]
[0]		[N-1]	[0]		[N-1]		[0]		[N-1]

# Example Exam Question

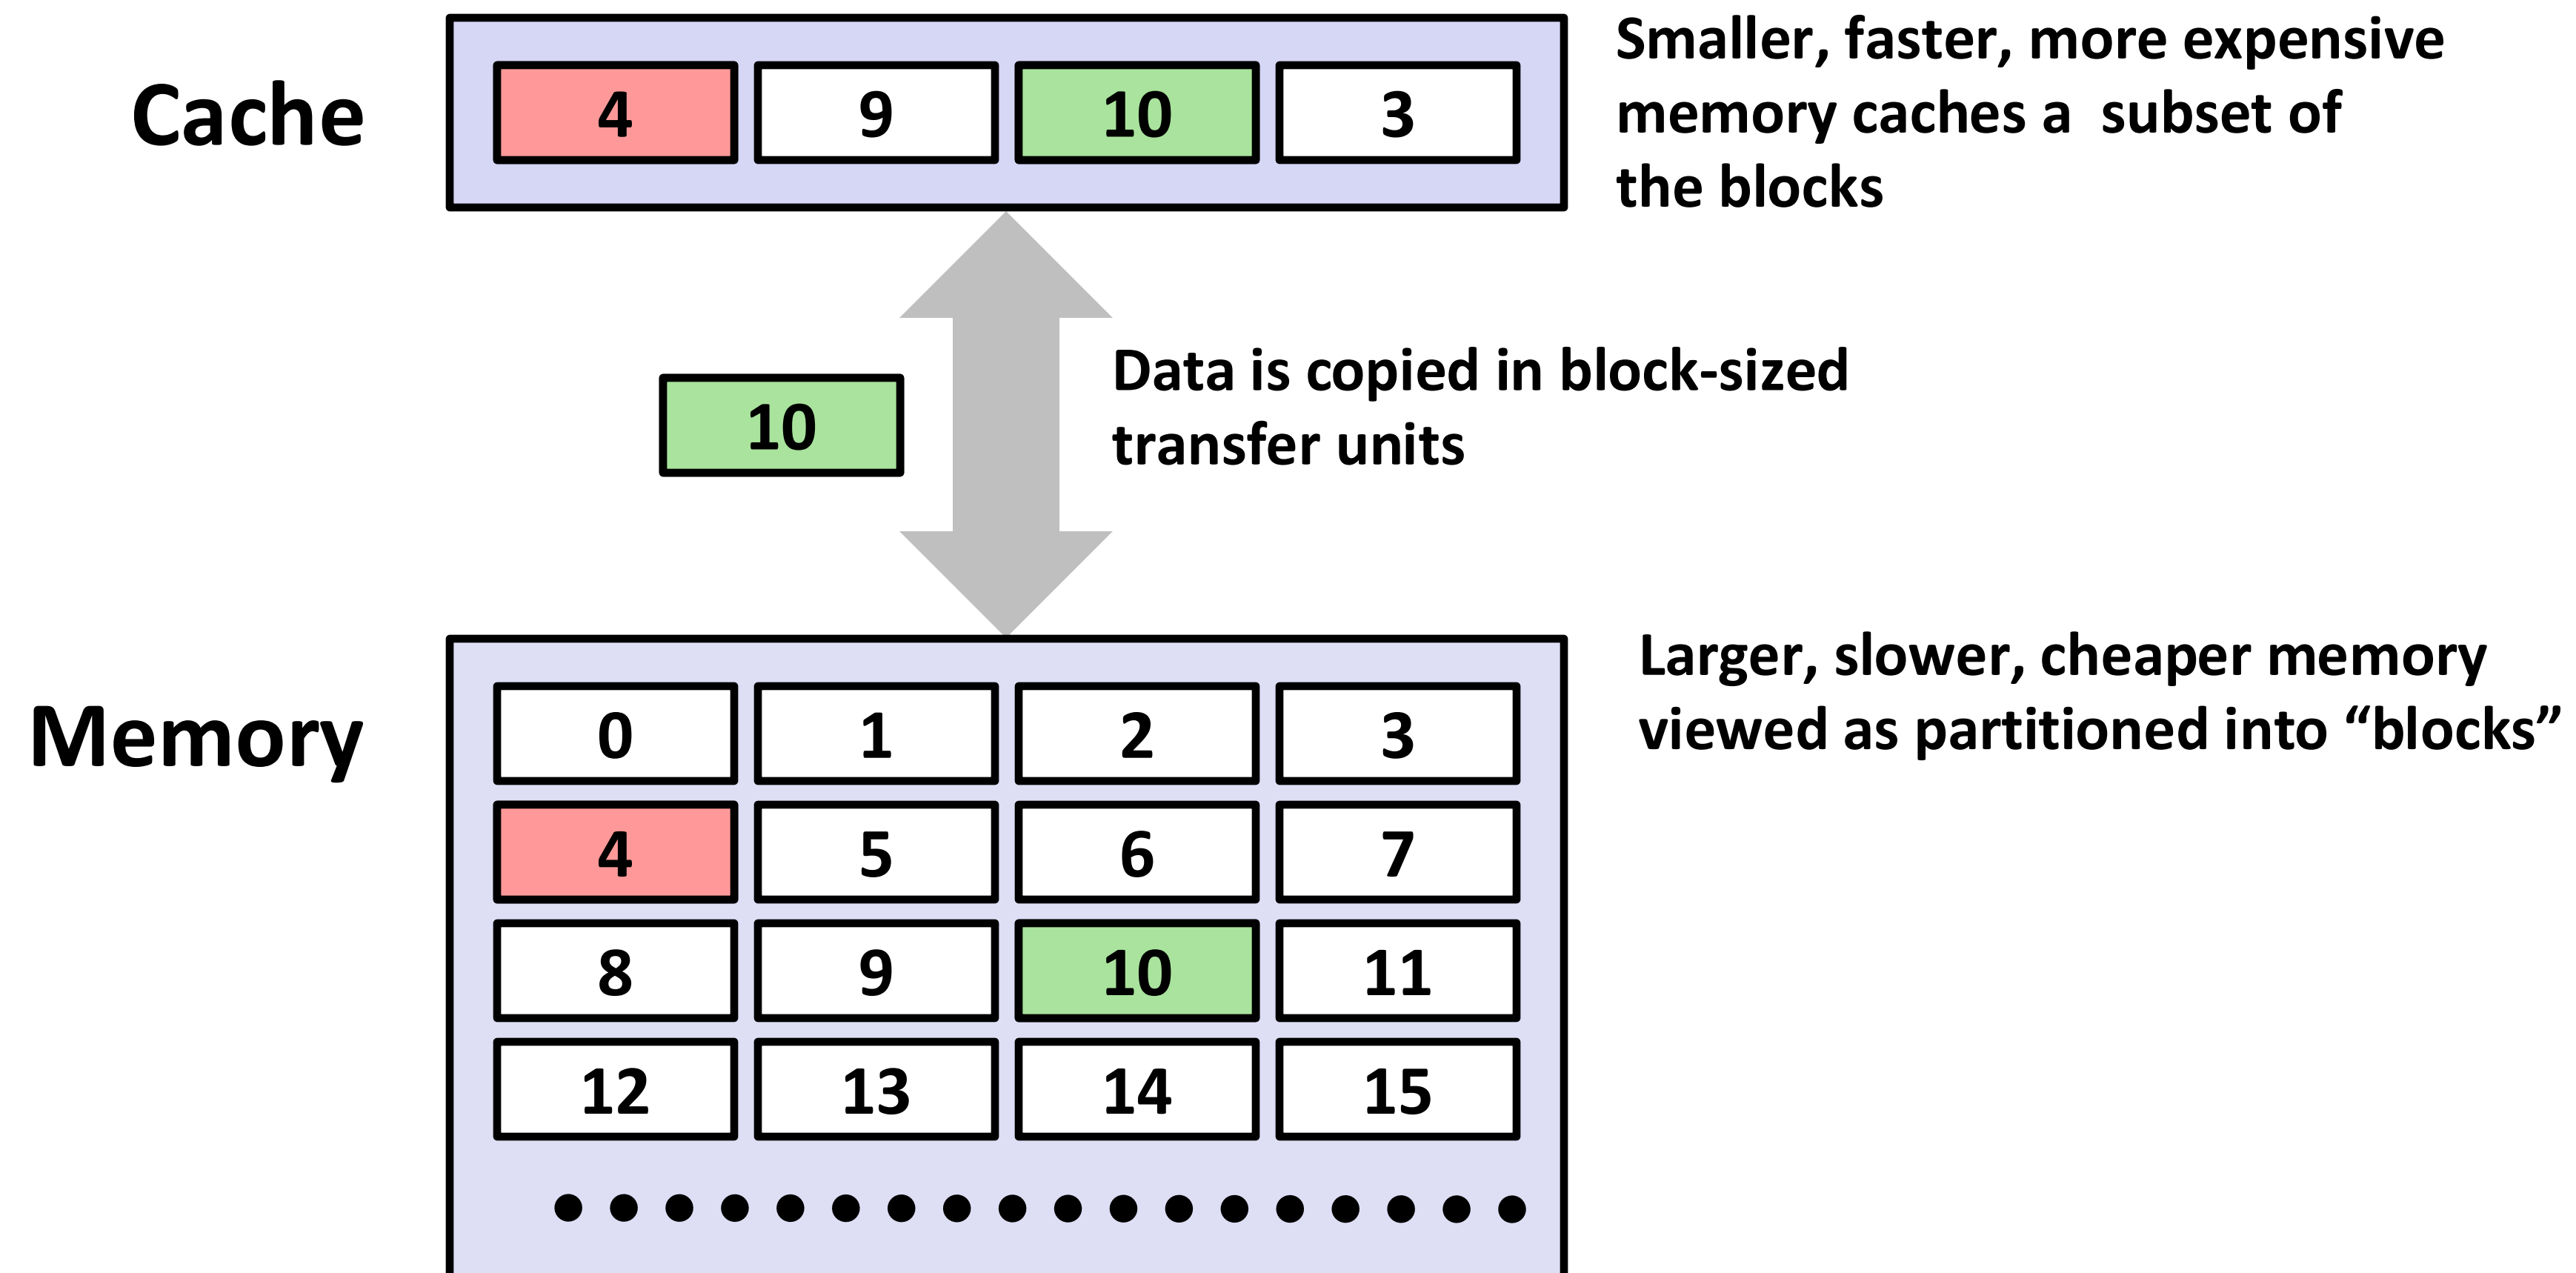
```
int sum_array_3d(int a[M][N][N])
{
    int i, j, k, sum = 0;

    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            for (k = 0; k < M; k++)
                sum += a[k][i][j];


    return sum;
}
```

- **Question:** Can you permute the loops so that the function scans the 3-d array *a* with a stride-1 reference pattern (and thus has good spatial locality)?

# Cache in action



# General Cache Concepts: Hit




*Data in block 14 is needed*

*Block 14 is in cache:*  
***Hit!***



# General Cache Concepts: Miss



*Data in block 12 is needed*

*Block 12 is not in cache:*  
**Miss!**

*Block 12 is fetched from  
memory*

*Block 12 is stored in cache*

- **Placement policy:**  
determines where b goes
- **Replacement policy:**  
determines which block  
gets evicted (victim)


# Cache in action

- If always cache hit, bandwidth?
- If always cache miss, bandwidth?

**Processor**

**Cache**

**Memory**




# Open Question in Cache: ChatGPT

- ChatGPT: every time ChatGPT outputs token, it needs to see 350 GB parameters
- How to optimize this?

**Processor**

**Cache**

**Memory**



# Foundation of Data Systems

- Computer Organization
  - Representation of data
  - processors, memory, storage
- **OS basics**
  - Process, scheduling
  - Memory

# What is Operation System?

- Layers between applications and hardware



- OS makes computer hardware useful to programmers
  - Otherwise, users need to speak machine code to computer
- **[Usually]** Provides abstractions for applications
  - Manages and hides details of hardware
  - Accesses hardware through low/level interfaces unavailable to applications
- **[Often]** Provides protection
  - Prevents one app/user from clobbering another

# A Primitive OS v1

- OS v1: just a library of standard services [no protection]



OS: interfaces above hw drivers

Hardware

- Simplifying assumptions:
  - System runs one program at a time
  - No bad users or programs (?)
- Problem: poor utilization
  - poor utilization of hardware (e.g., CPU idle while waiting for disk)
  - poor utilization of human user (must wait for each program to finish)

# OS v2: Multi-tasking

- Say: we extend the OS a bit to support many APPs
  - When one process blocks (waiting for disk, network, user input, etc.) run another process



- Problem: What can ill-behaved process do?
  - Go into infinite loop and never relinquish CPU
  - Scribble over other processes' memory to make them fail
- OS provides mechanisms **protection** to address these problems:
  - Preemption – take CPU away from looping process
  - Memory protection – protect one process' memory from one another

# What is A Real OS?

- OS: manage and assign hardware resources to apps
- Goal: with N users/apps, system not N times slower
  - **Idea:** Giving resources to users who actually need them
- What can go wrong?
  - One app can interfere with other app (need **isolation**)
  - Users are gluttons, use too much CPU, etc. (need **scheduling**)
  - Total memory usage of all apps/users greater than machine's RAM  
(need **memory management**)
  - Disks are shared across apps / users and must be arranged properly  
(need **file systems**)



# Summary of OS: a software between apps and hardware

- Goal 1: Provide convenience to users
- Goal 2: Efficiency -- Manage compute, memory, storage resources
  - Goal 2.1: Running N processes Not N times slower
    - As fast as possible
  - Goal 2.2: Running N apps
    - Even when their total memory >> physical memory cap
- Goal 3: Provide **protection**
  - One process won't mess up the entire computer
  - One process won't mess up with other processes

Process management

Memory management

System calls

# Summary of OS: a software between apps and hardware

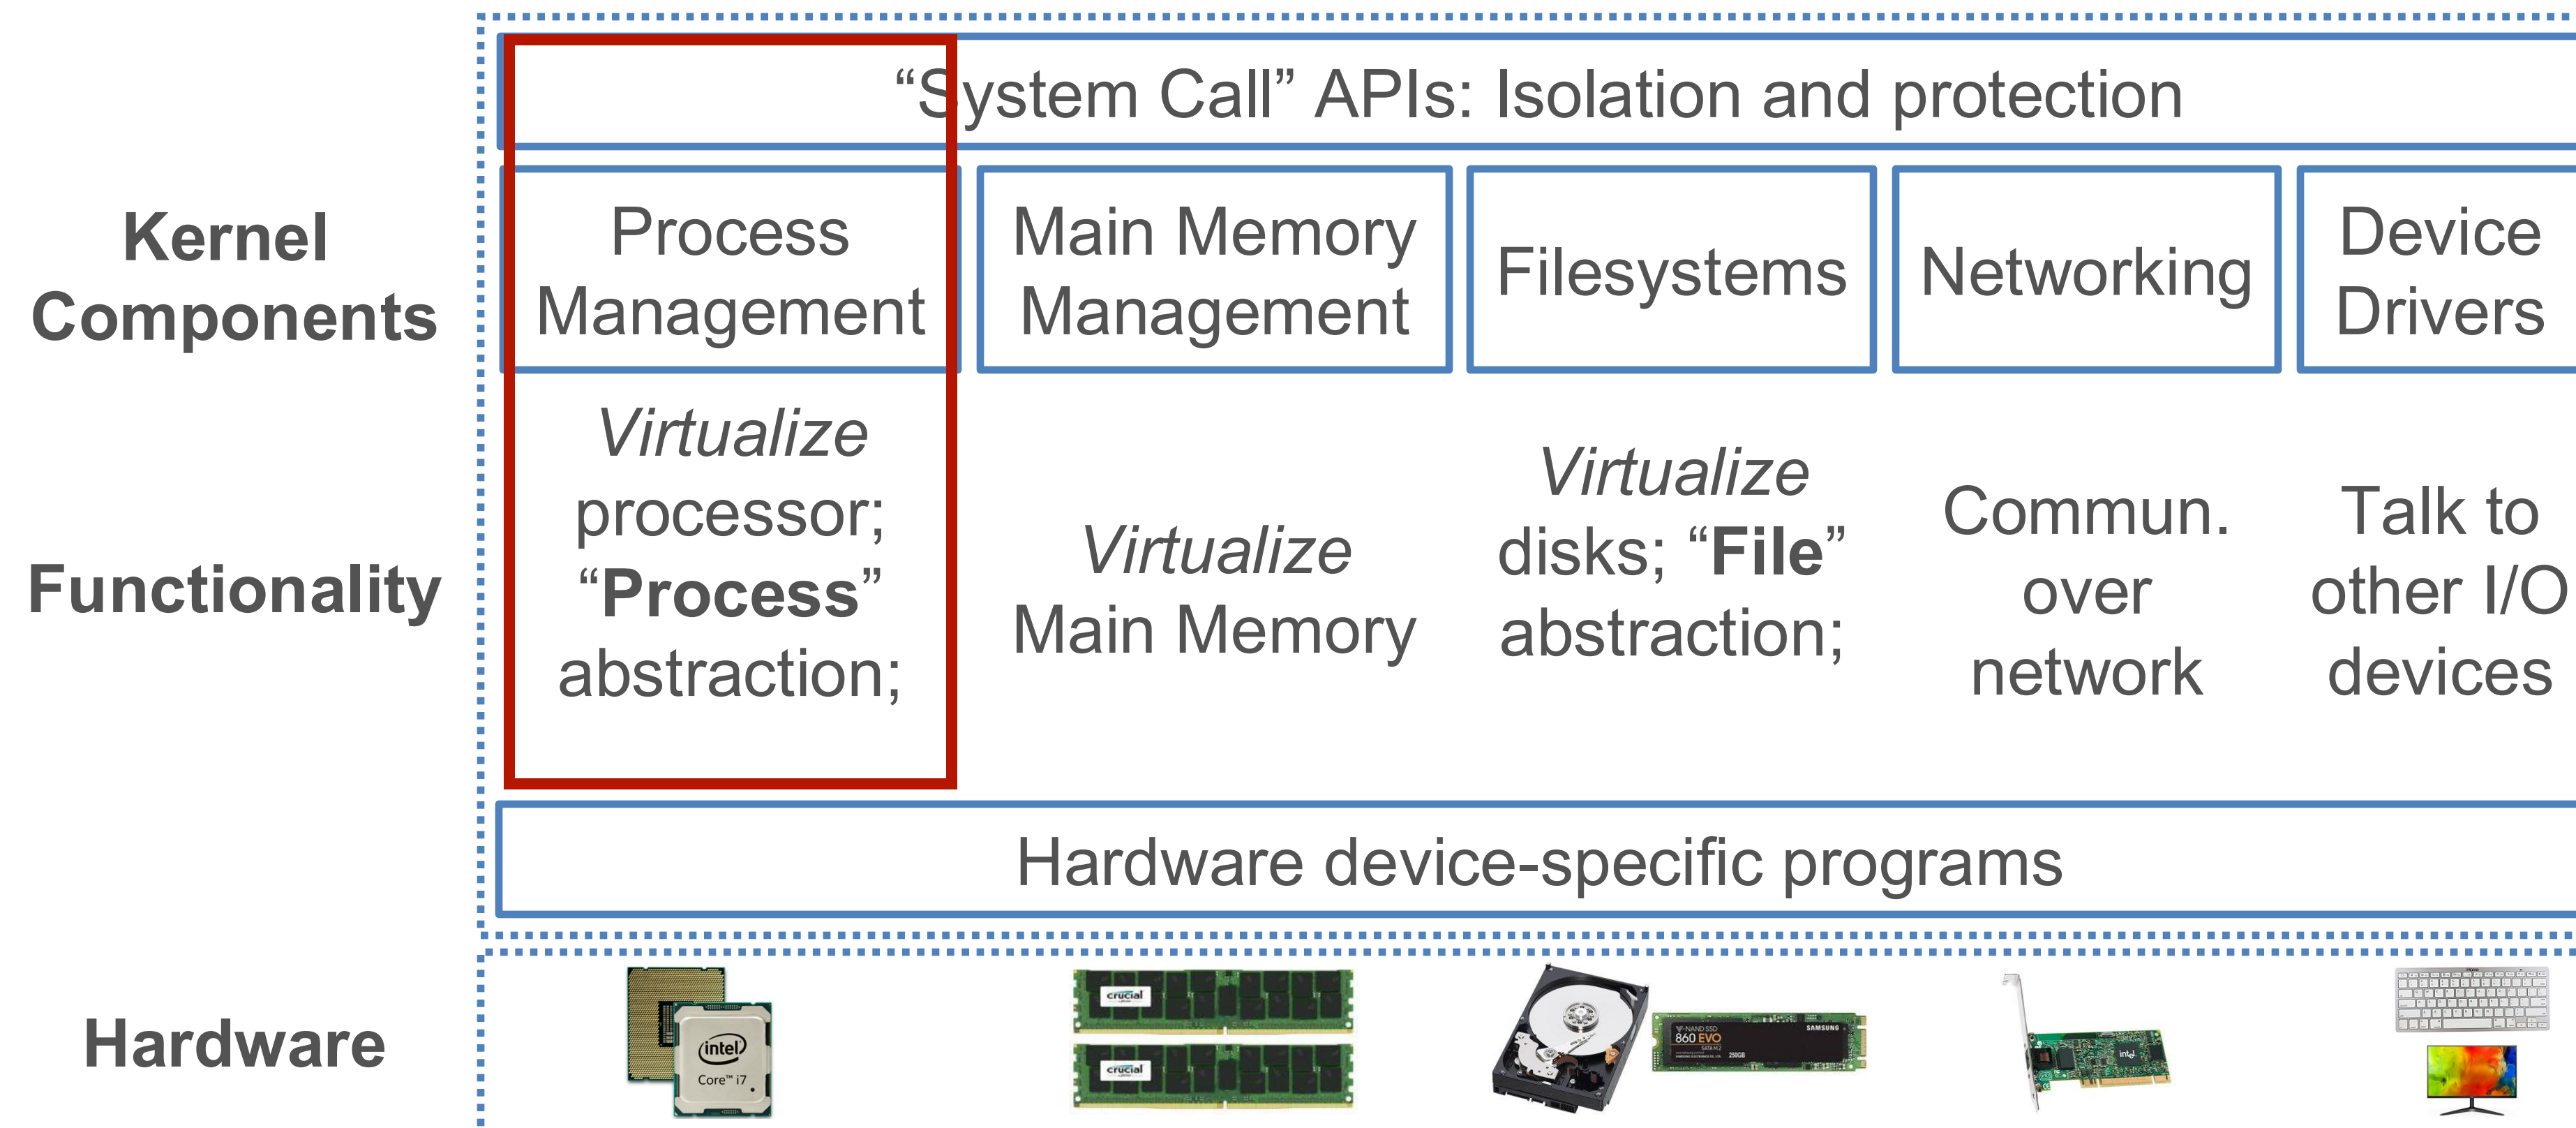
- Goal 1: Provide convenience to users
- Goal 2: Efficiency -- Manage compute, memory, storage resources
  - Goal 2.1: Running N processes Not N times slower
    - As fast as possible
  - Goal 2.2: Running N apps
    - Even when their total memory  $\gg$  physical memory cap
- **Goal 3: Provide protection**
  - **One process won't mess up the entire computer**
  - **One process won't mess up with other processes**

Process management  
Memory management

System calls

# OS provides Isolation using System Calls

- **System call:** The layer for isolation -- it abstracts the hardware and APIs for programs to use



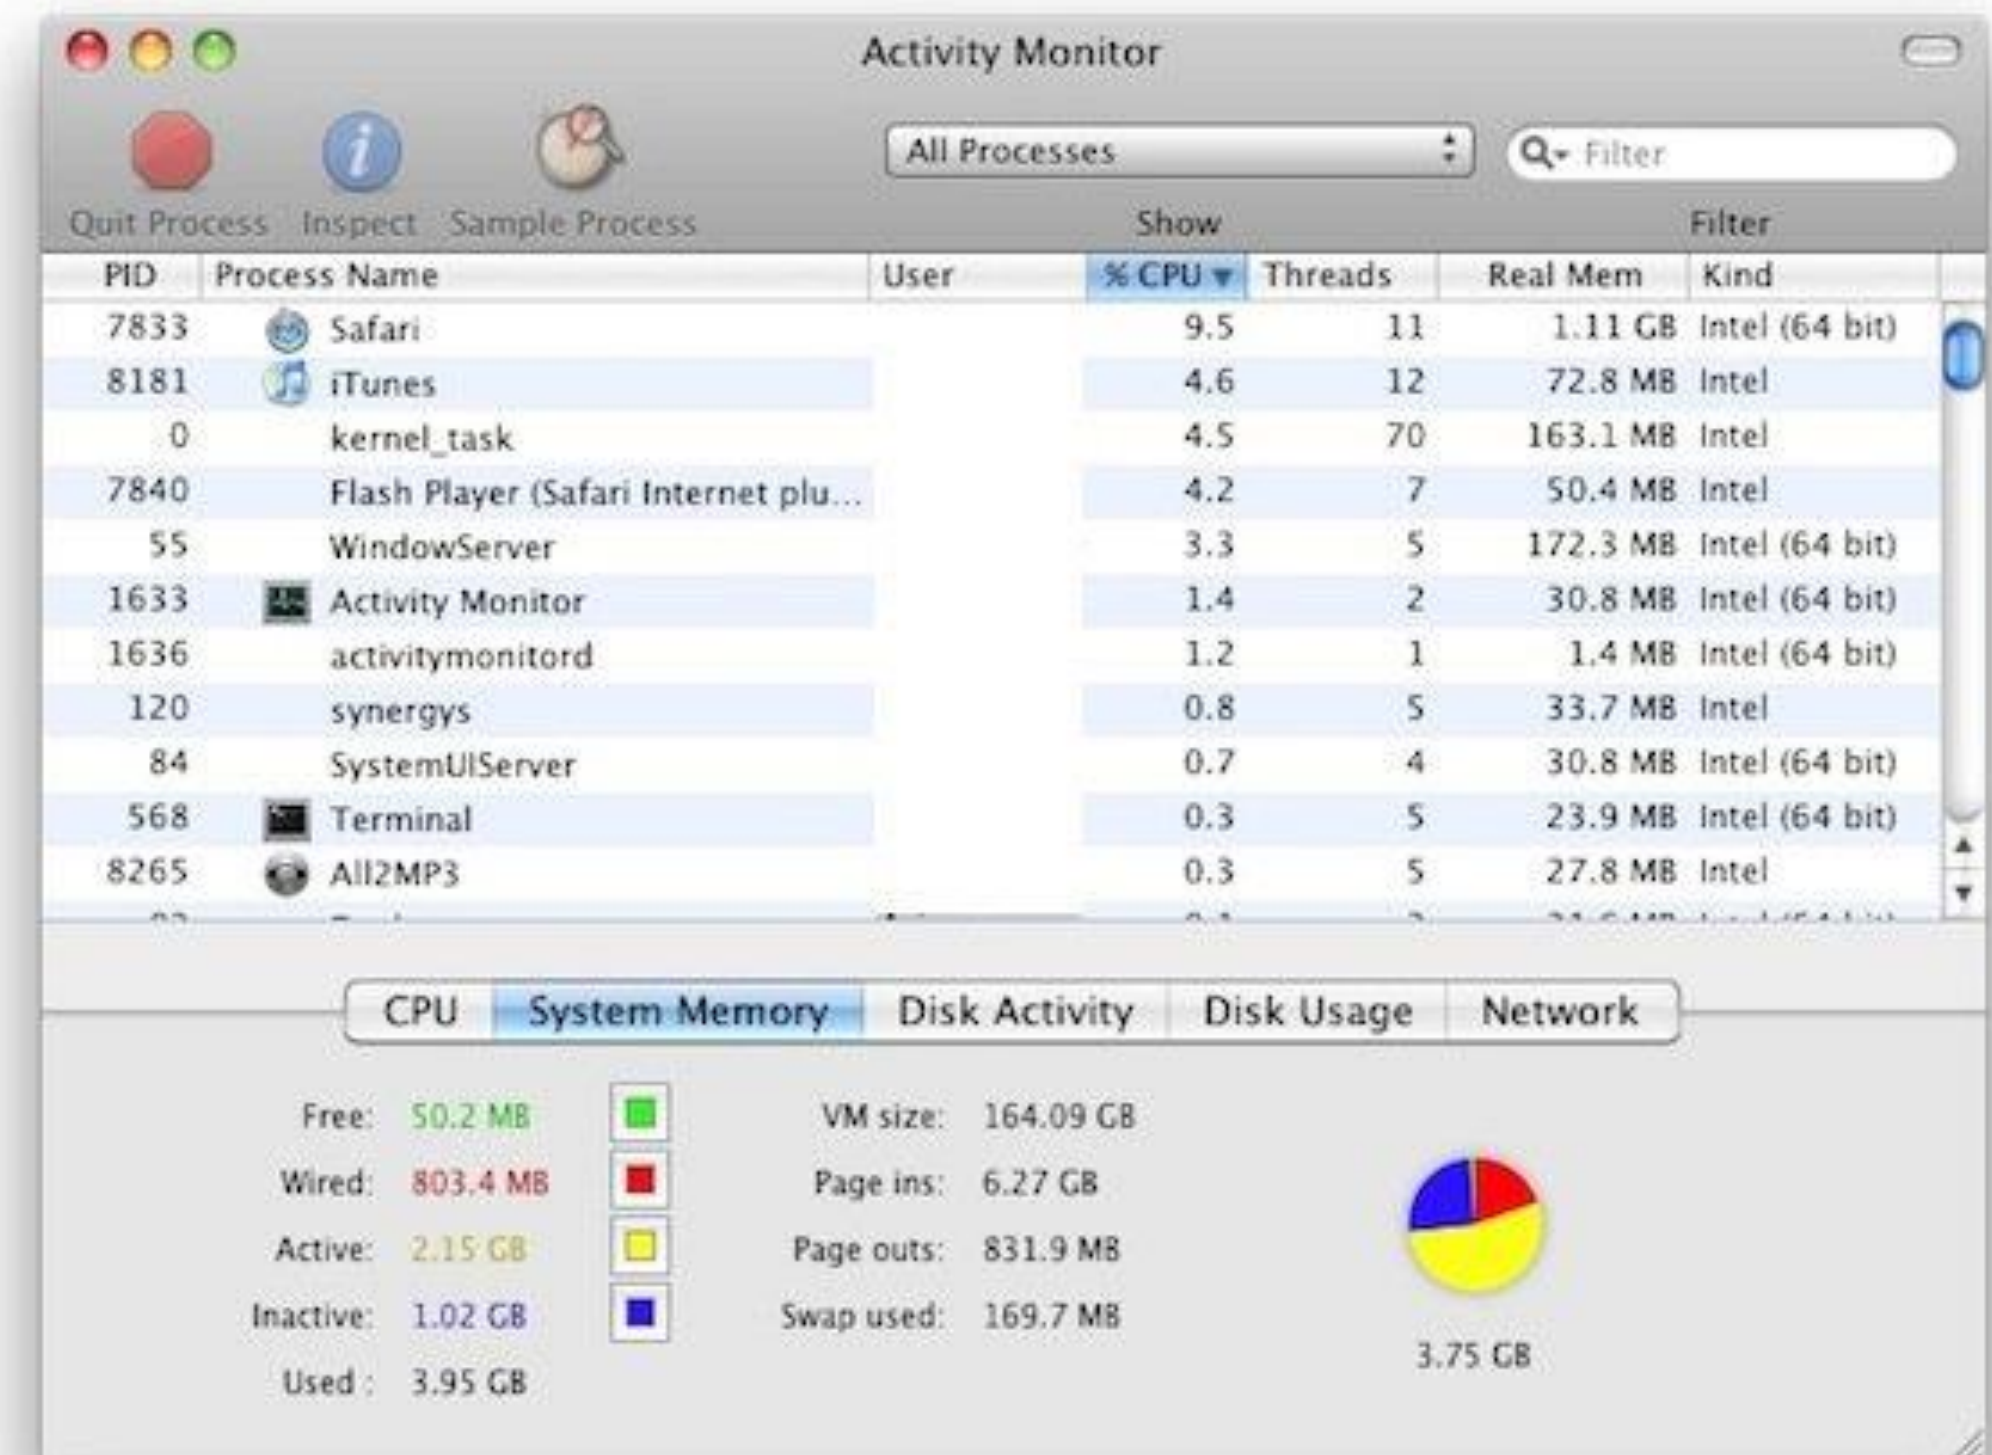
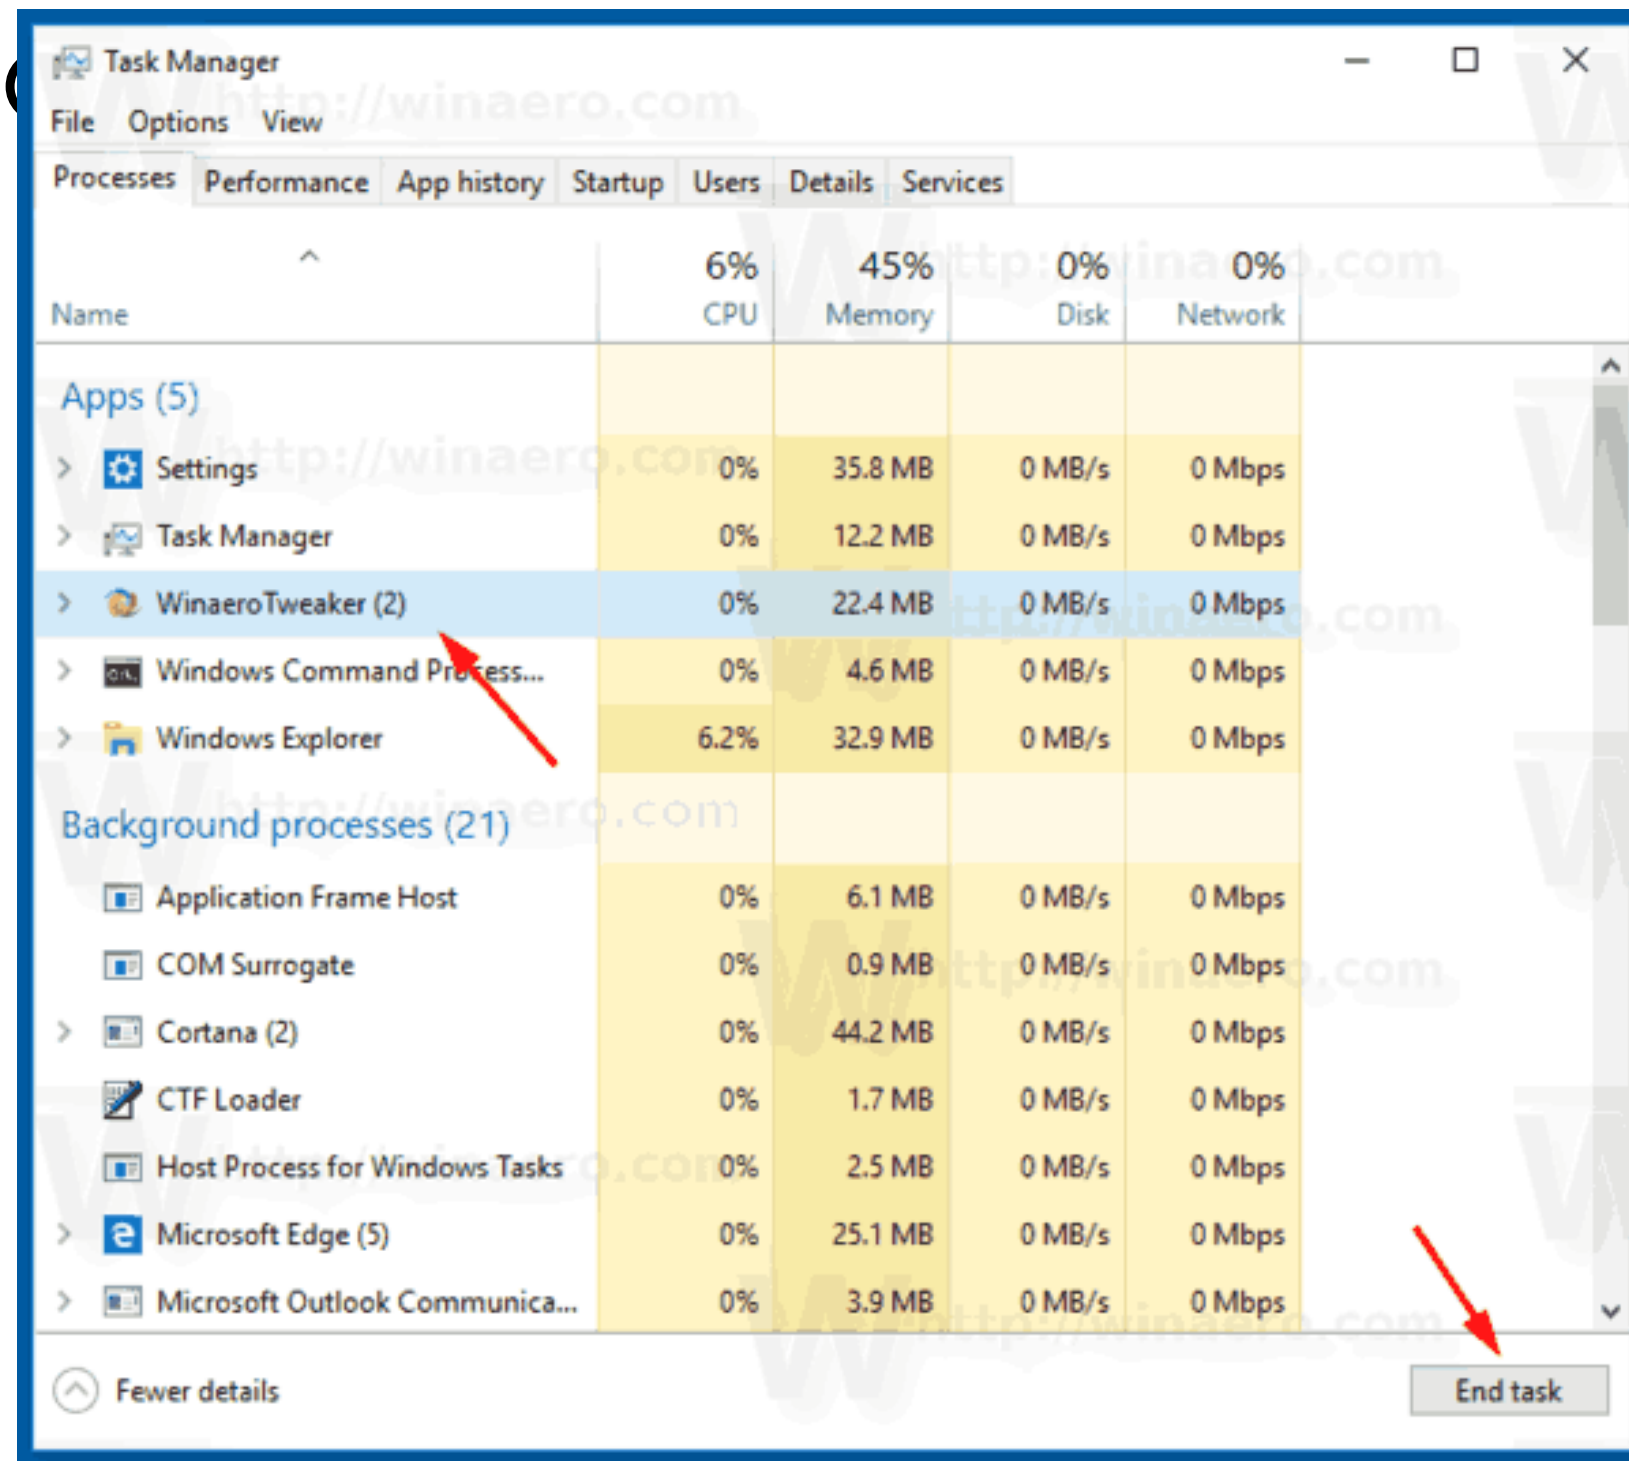
# Foundation of Data Systems

- Computer Organization
  - Representation of data
  - processors, memory, storage
- OS basics
  - **Process, scheduling**
  - Memory

# Processes - the central abstraction in OS

- Definition: A *process* is an instance of a running program.
- One of the most profound ideas in computer science


- No or



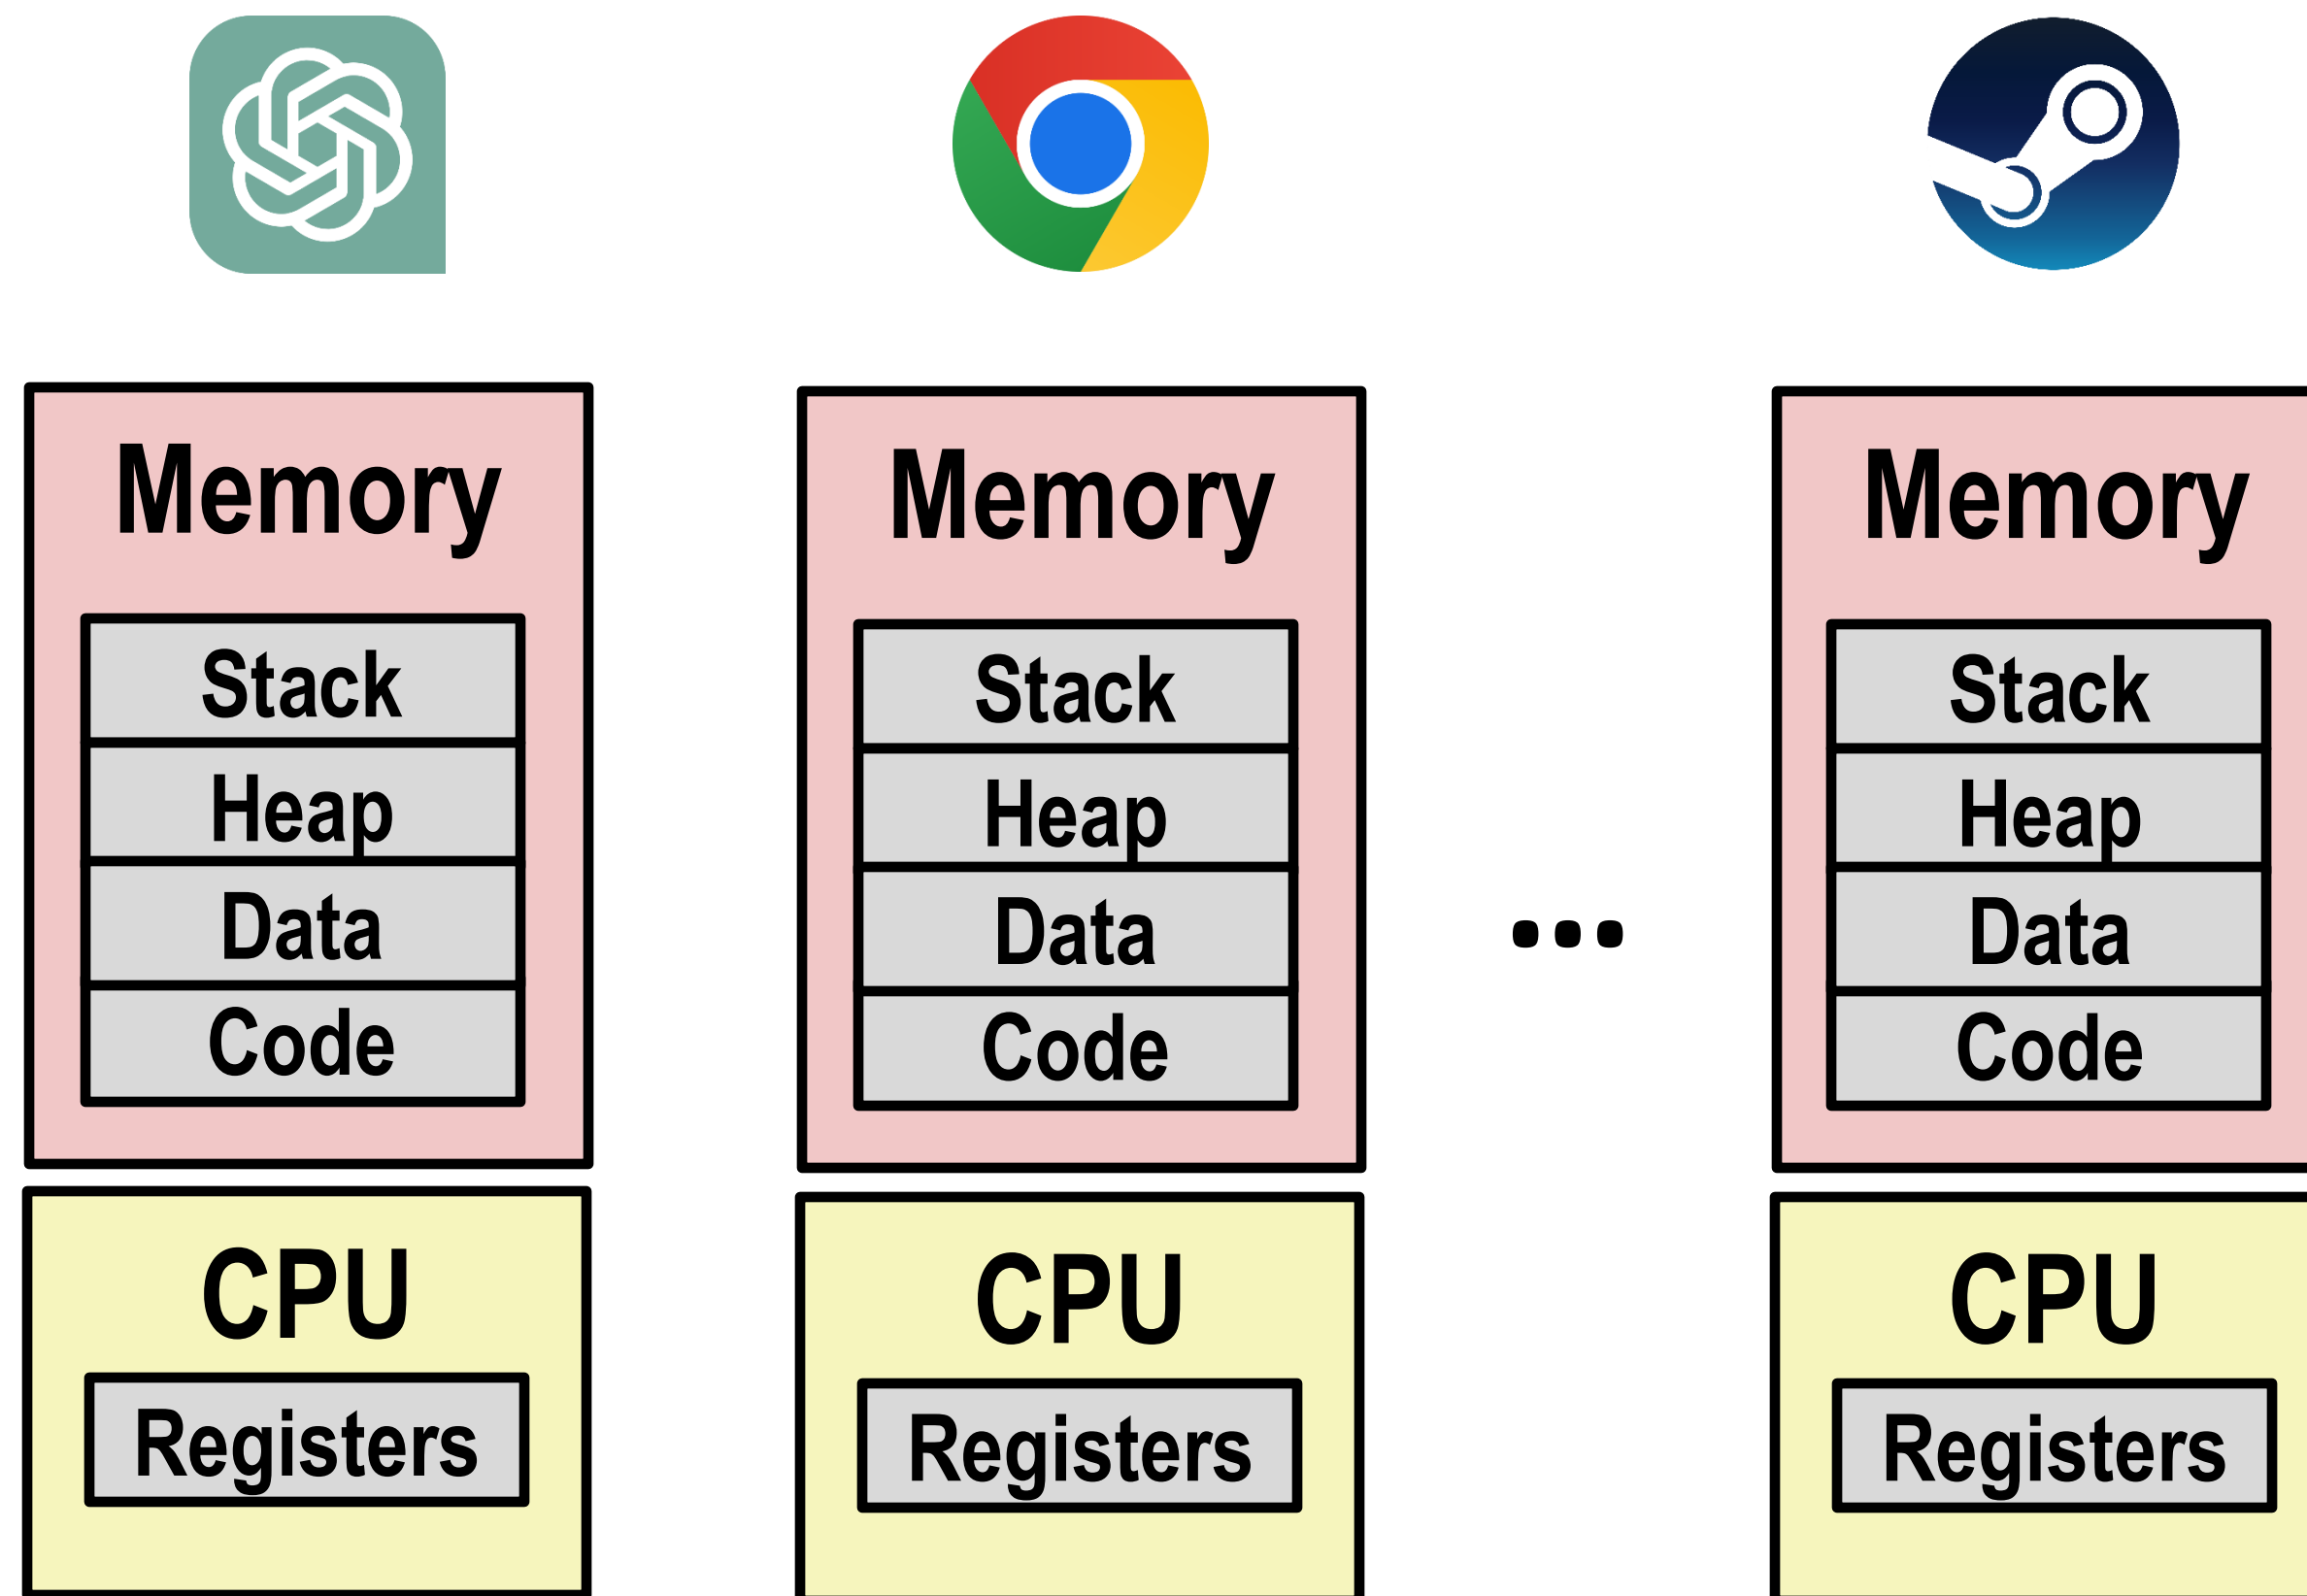


# Processes - the central abstraction in OS

- Process provides each program with two key abstractions (for resources):
  - **Compute Resource**
    - Each program seems to have exclusive use of the CPU
    - Provided by kernel mechanism called *context switching*
  - **Memory Resource**
    - Each program seems to have exclusive use of main mem
    - Provided by kernel mechanism called virtual memory



# Multiprocessing in OS: The Illusion



- Computer runs many processes simultaneously

# Multiprocessing Example

top command in terminal: many processes, Identified by Process ID (**PID**)

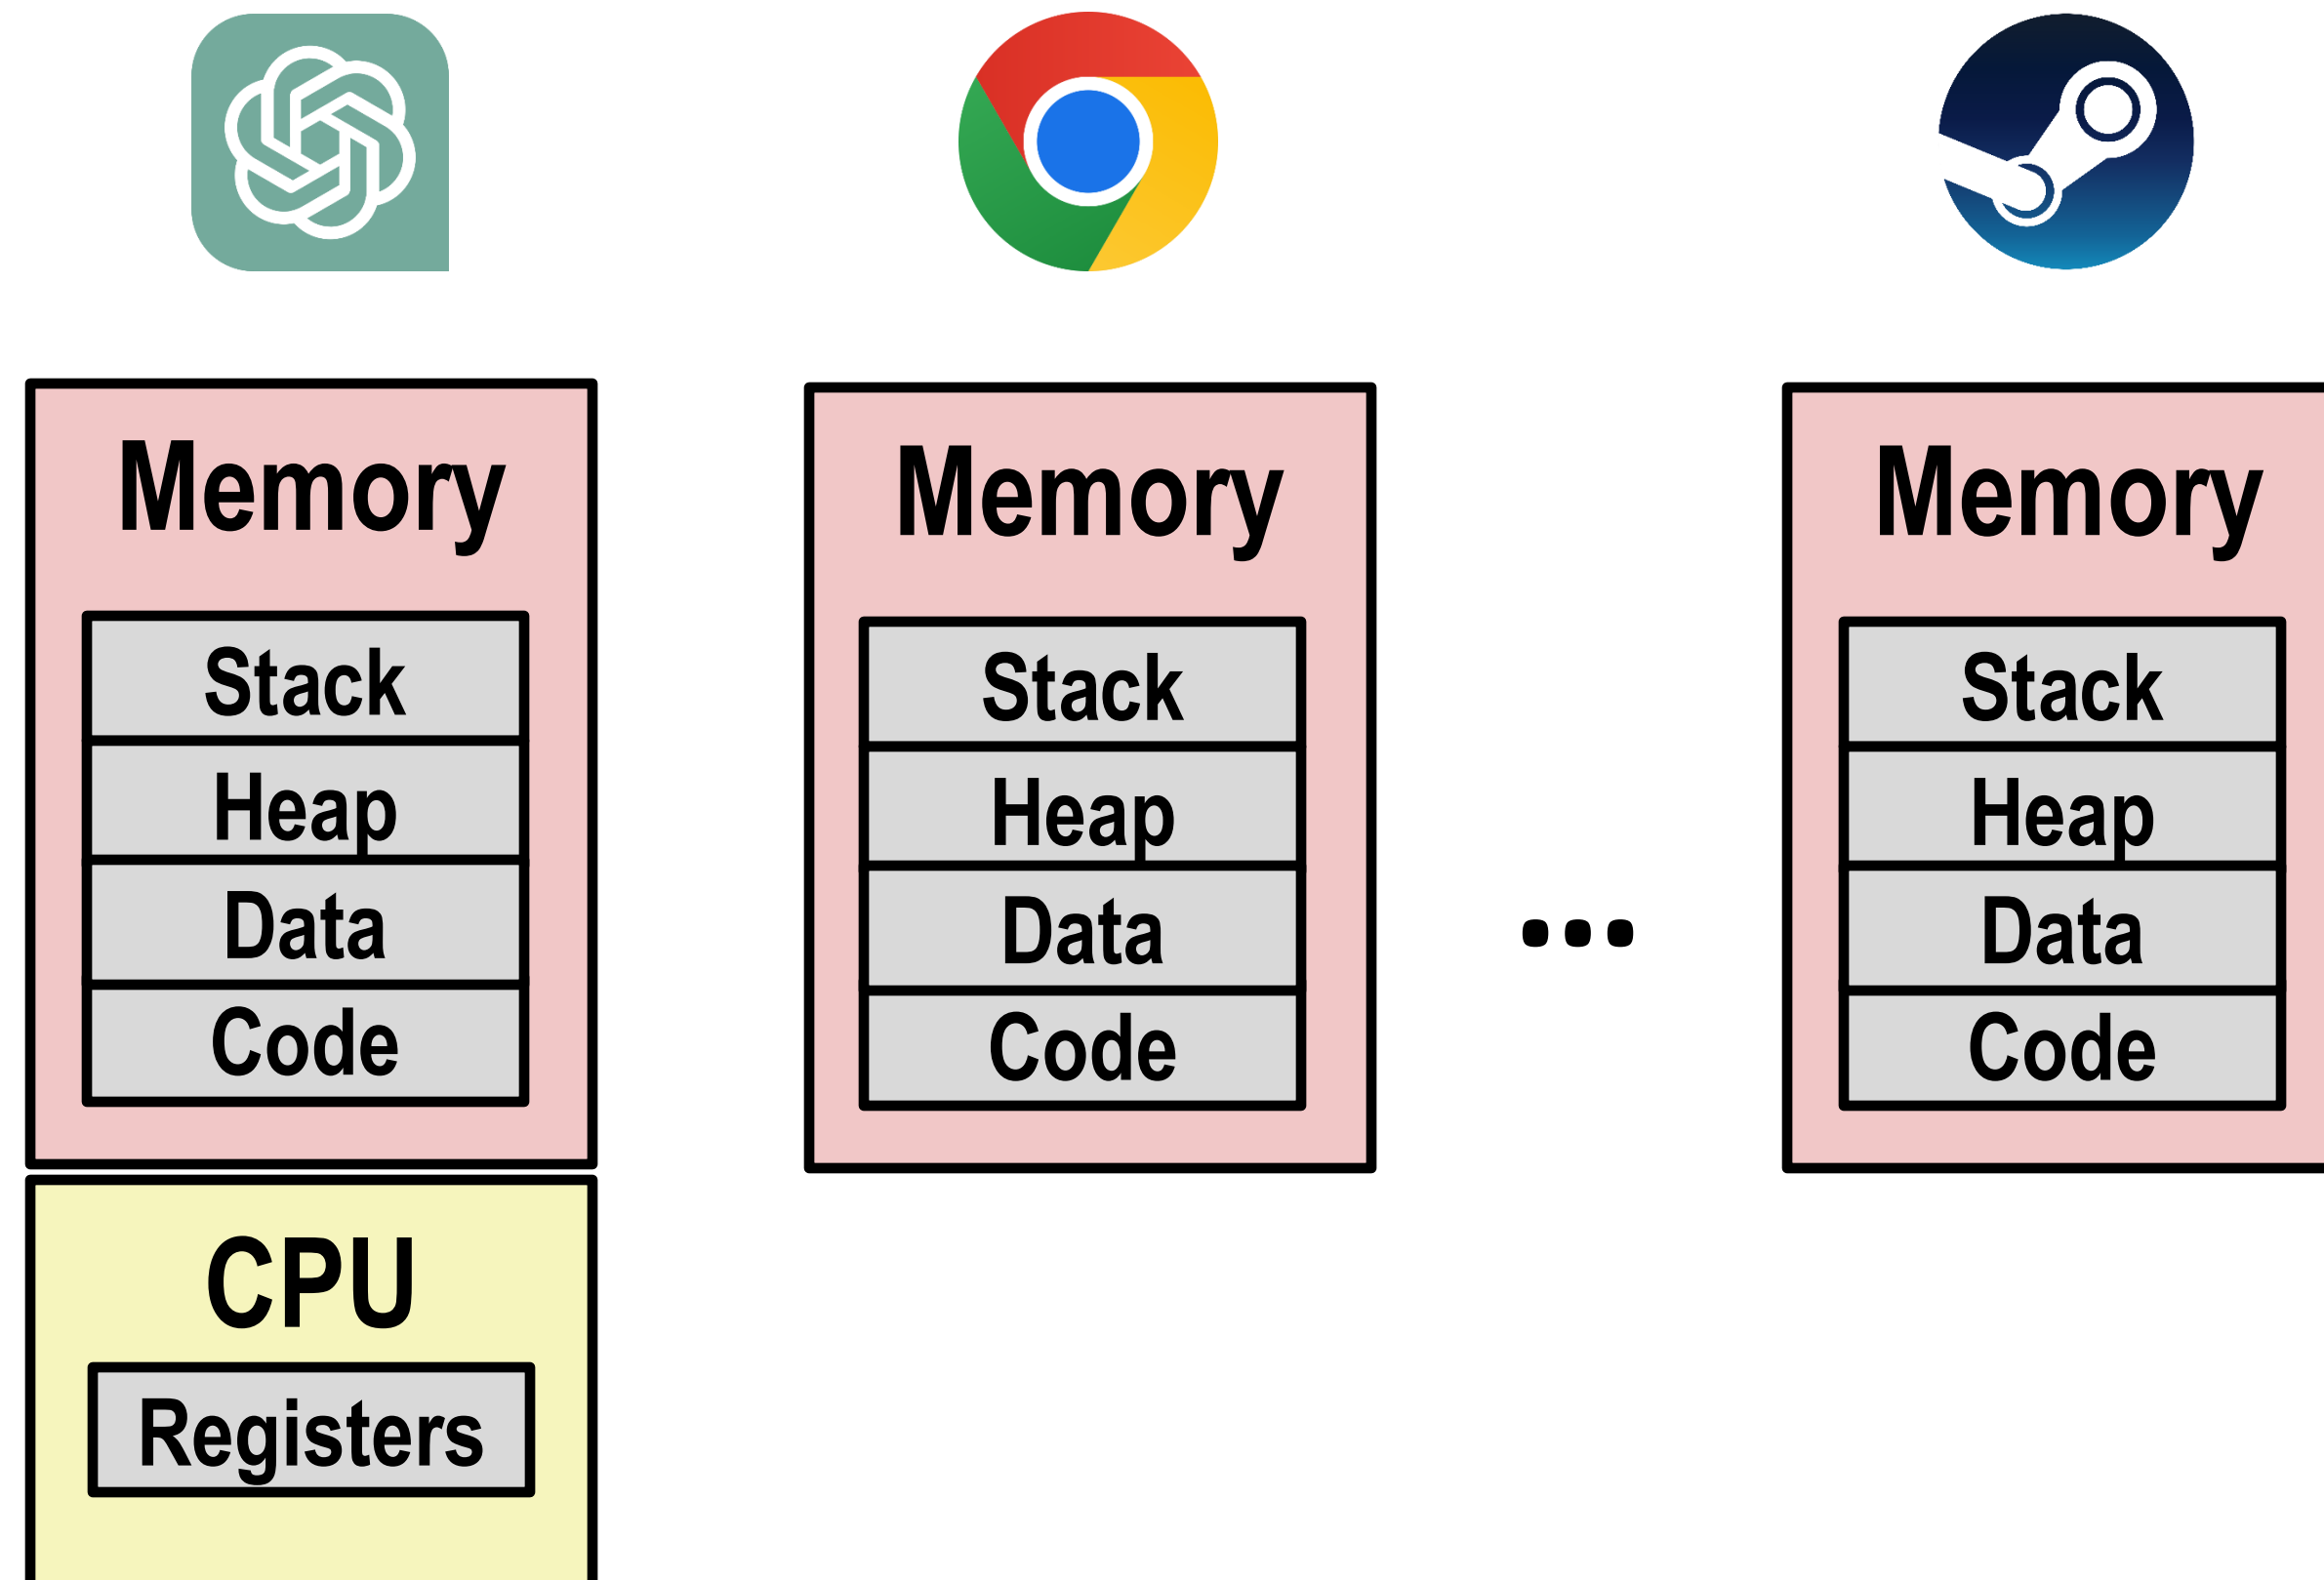
```
Processes: 123 total, 5 running, 9 stuck, 109 sleeping, 611 threads
Load Avg: 1.03, 1.13, 1.14  CPU usage: 3.27% user, 5.15% sys, 91.56% idle
SharedLibs: 576K resident, 0B data, 0B linkedit.
MemRegions: 27958 total, 1127M resident, 35M private, 494M shared.
PhysMem: 1039M wired, 1974M active, 1062M inactive, 4076M used, 18M free.
VM: 280G vsize, 1091M framework vsize, 23075213(1) pageins, 5843367(0) pageouts.
Networks: packets: 41046228/11G in, 66083096/77G out.
Disks: 17874391/349G read, 12847373/594G written.

PID      COMMAND     %CPU TIME    #TH   #WQ   #PORT #MREG RPRVT  RSHRD  RSIZE  VPRVT  VSIZE
99217- Microsoft Of 0.0 02:28.34 4     1    202   418   21M   24M   21M   66M   763M
99051  usbmuxd     0.0 00:04.10 3     1     47    66   436K  216K  480K  60M   2422M
99006  iTunesHelper 0.0 00:01.23 2     1     55    78   728K  3124K 1124K  43M   2429M
84286  bash        0.0 00:00.11 1     0     20    24   224K  732K  484K  17M   2378M
84285  xterm       0.0 00:00.83 1     0     32    73   656K  872K  692K  9728K 2382M
55939- Microsoft Ex 0.3 21:58.97 10    3    360   954   16M   65M   46M   114M  1057M
54751  sleep       0.0 00:00.00 1     0     17    20    92K   212K  360K  9632K 2370M
54739  launchdadd  0.0 00:00.00 2     1     33    50   488K  220K  1736K  48M   2409M
54737  top         6.5 00:02.53 1/1   0     30    29  1416K  216K  2124K  17M   2378M
54719  automountd  0.0 00:00.02 7     1     53    64   860K  216K  2184K  53M   2413M
54701  ocspd       0.0 00:00.05 4     1     61    54  1268K  2644K  3132K  50M   2426M
54661  Grab        0.6 00:02.75 6     3    222+  389+  15M+  26M+  40M+  75M+  2556M+
54659  cookied     0.0 00:00.15 2     1     40    61  3316K  224K  4088K  42M   2411M
53818  mdworker    0.0 00:01.67 4     1     52    91  7628K  7412K  16M   48M   2438M
50878  mdworker    0.0 00:11.17 3     1     53    91  2464K  6148K  9976K  44M   2434M
50410  xterm       0.0 00:00.13 1     0     32    73   280K  872K  532K  9700K 2382M
50078  emacs       0.0 00:06.70 1     0     20    35    52K   216K  88K   18M   2392M
```



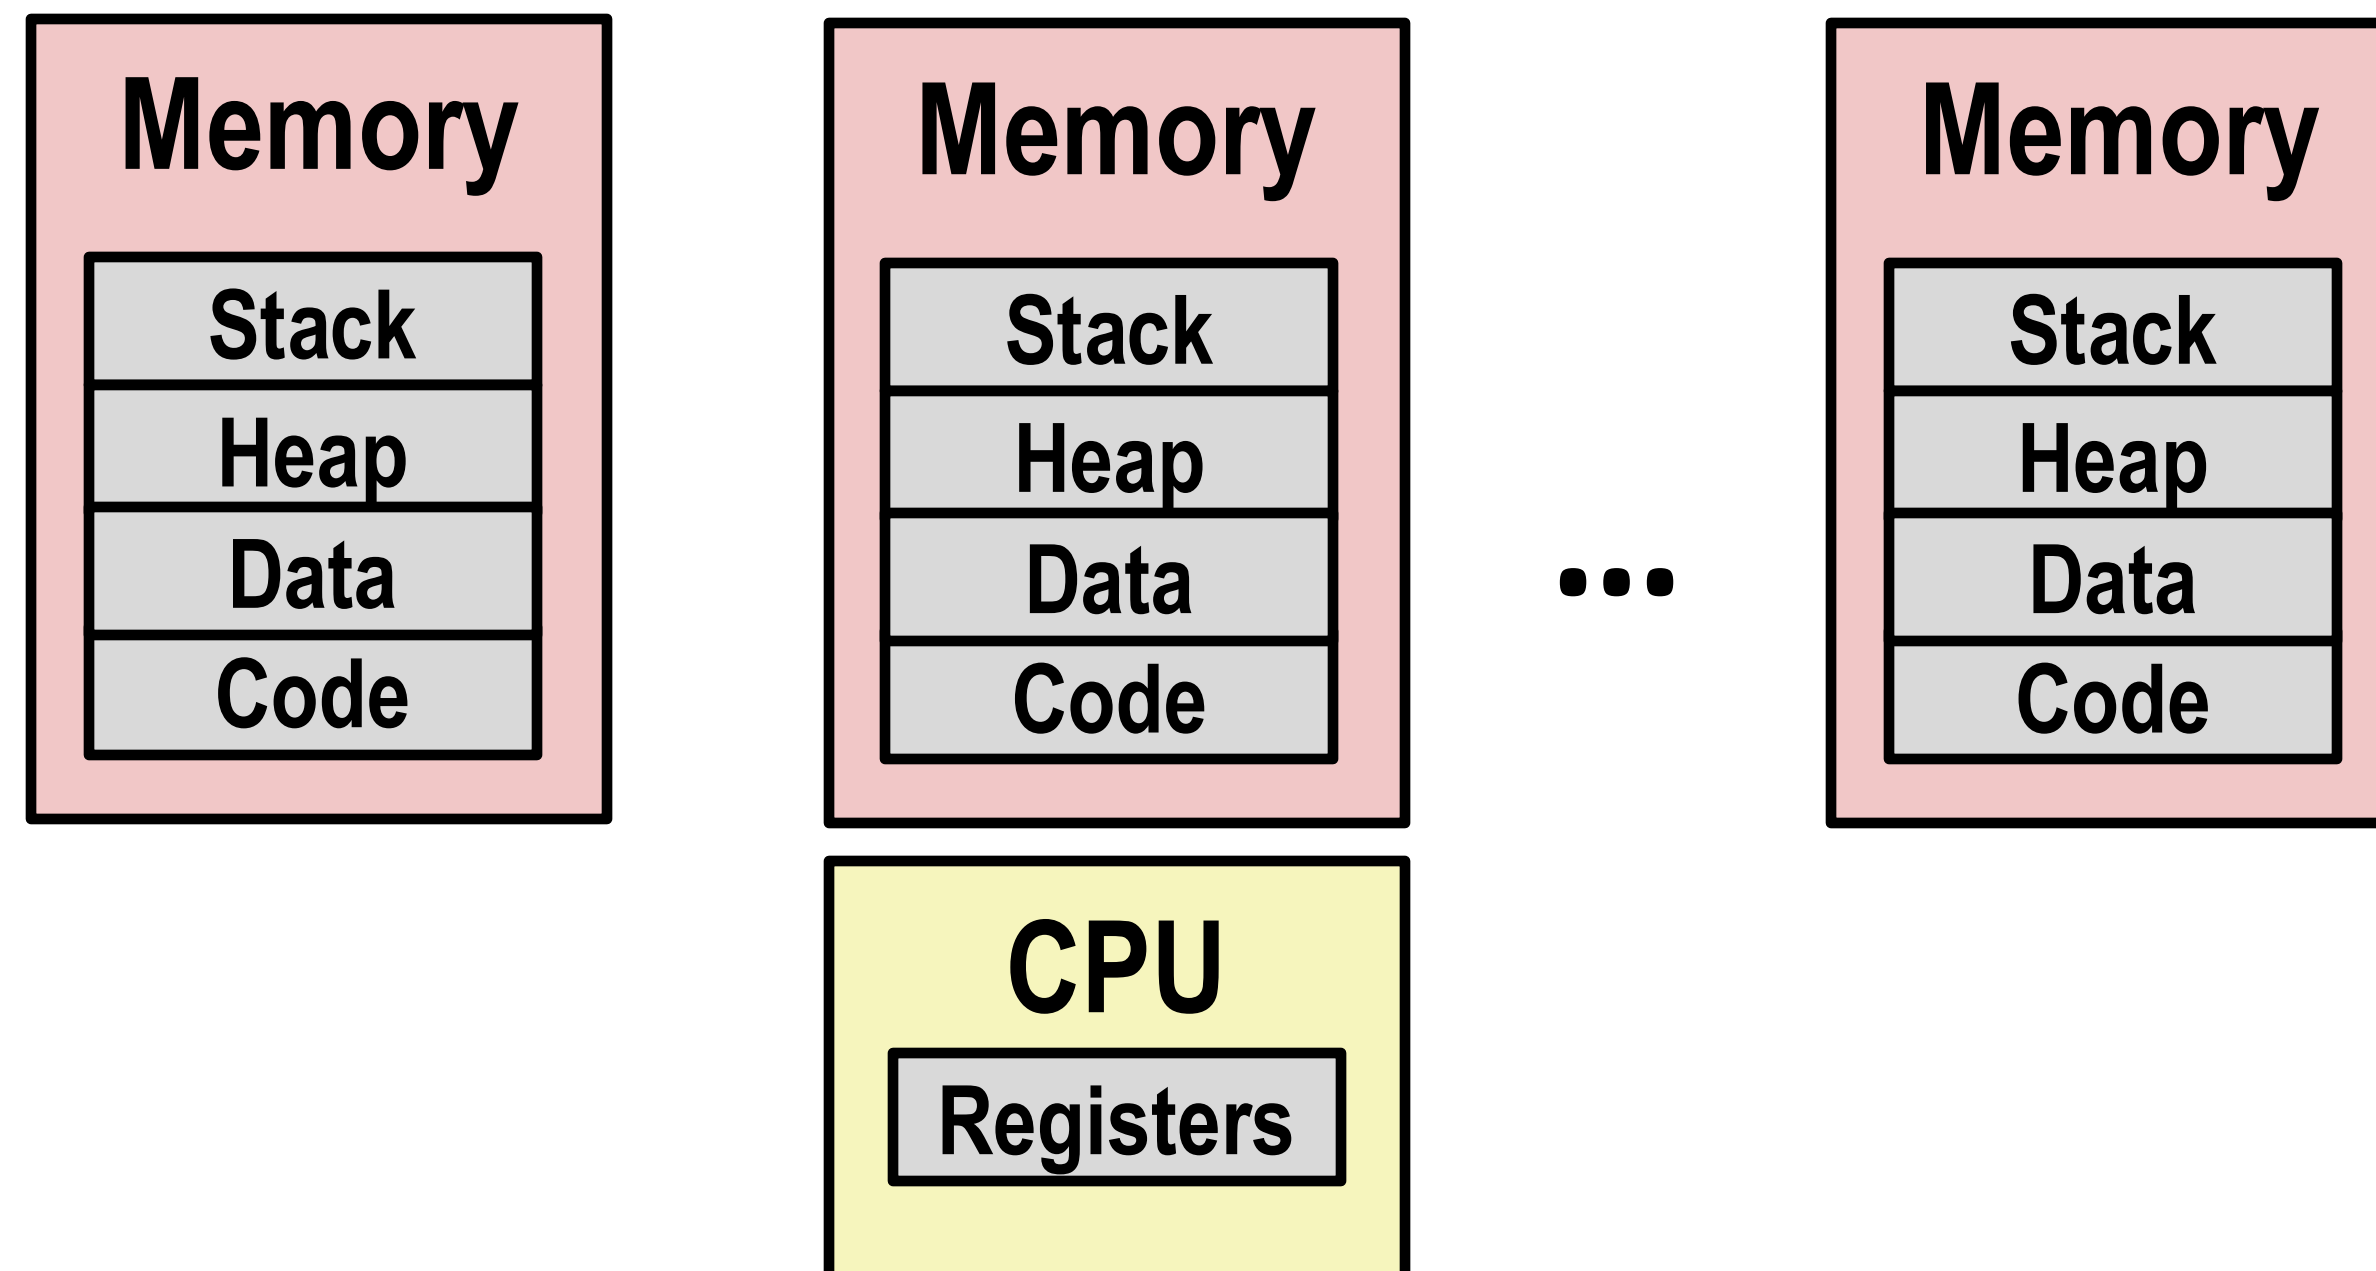
# Multiprocessing: A strawman solution

- Assign individual memory (say 1/3) to each APP
- Assign CPU to work on an APP until completion -> then next



# Multiprocessing: A strawman solution



- Assign individual memory (say 1/3) to each APP
- Assign CPU to work on an APP until completion -> then next





# Multiprocessing: A strawman solution

- Assign individual memory (say 1/3) to each APP
- Assign CPU to work on an APP until completion -> then

next



...



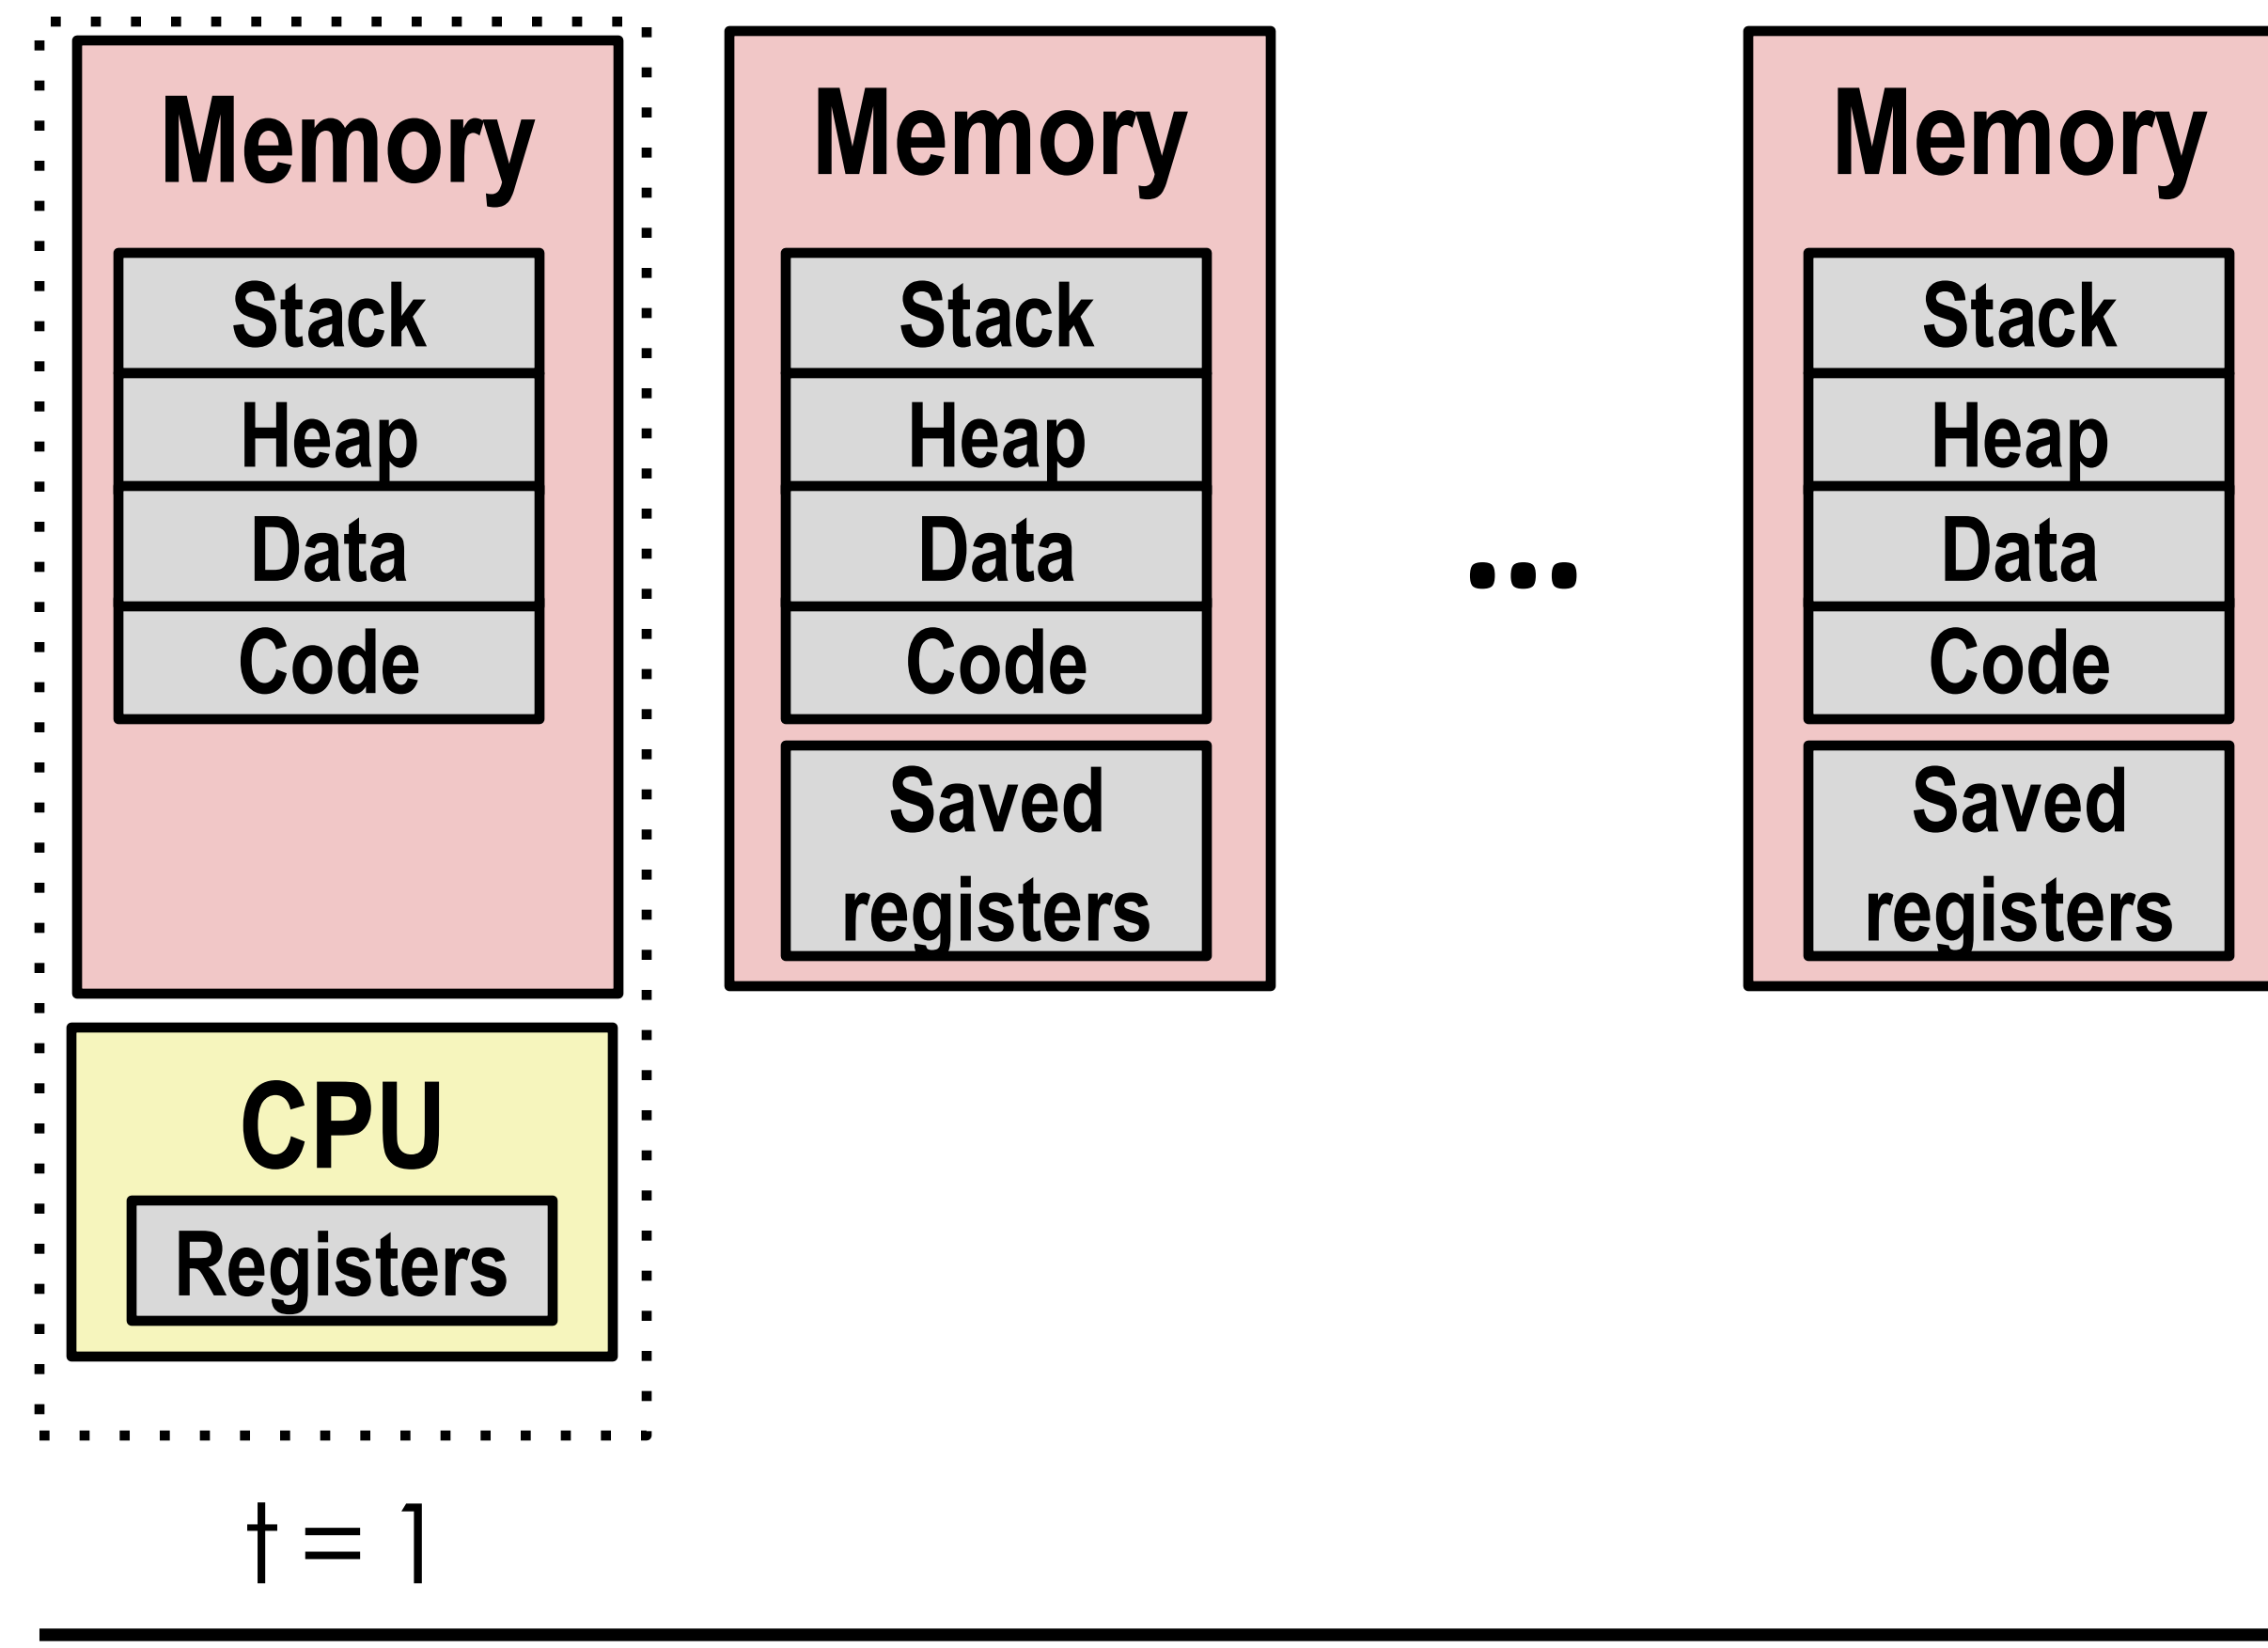
G1. Convenient?

G3: protection?

G2. Efficient?

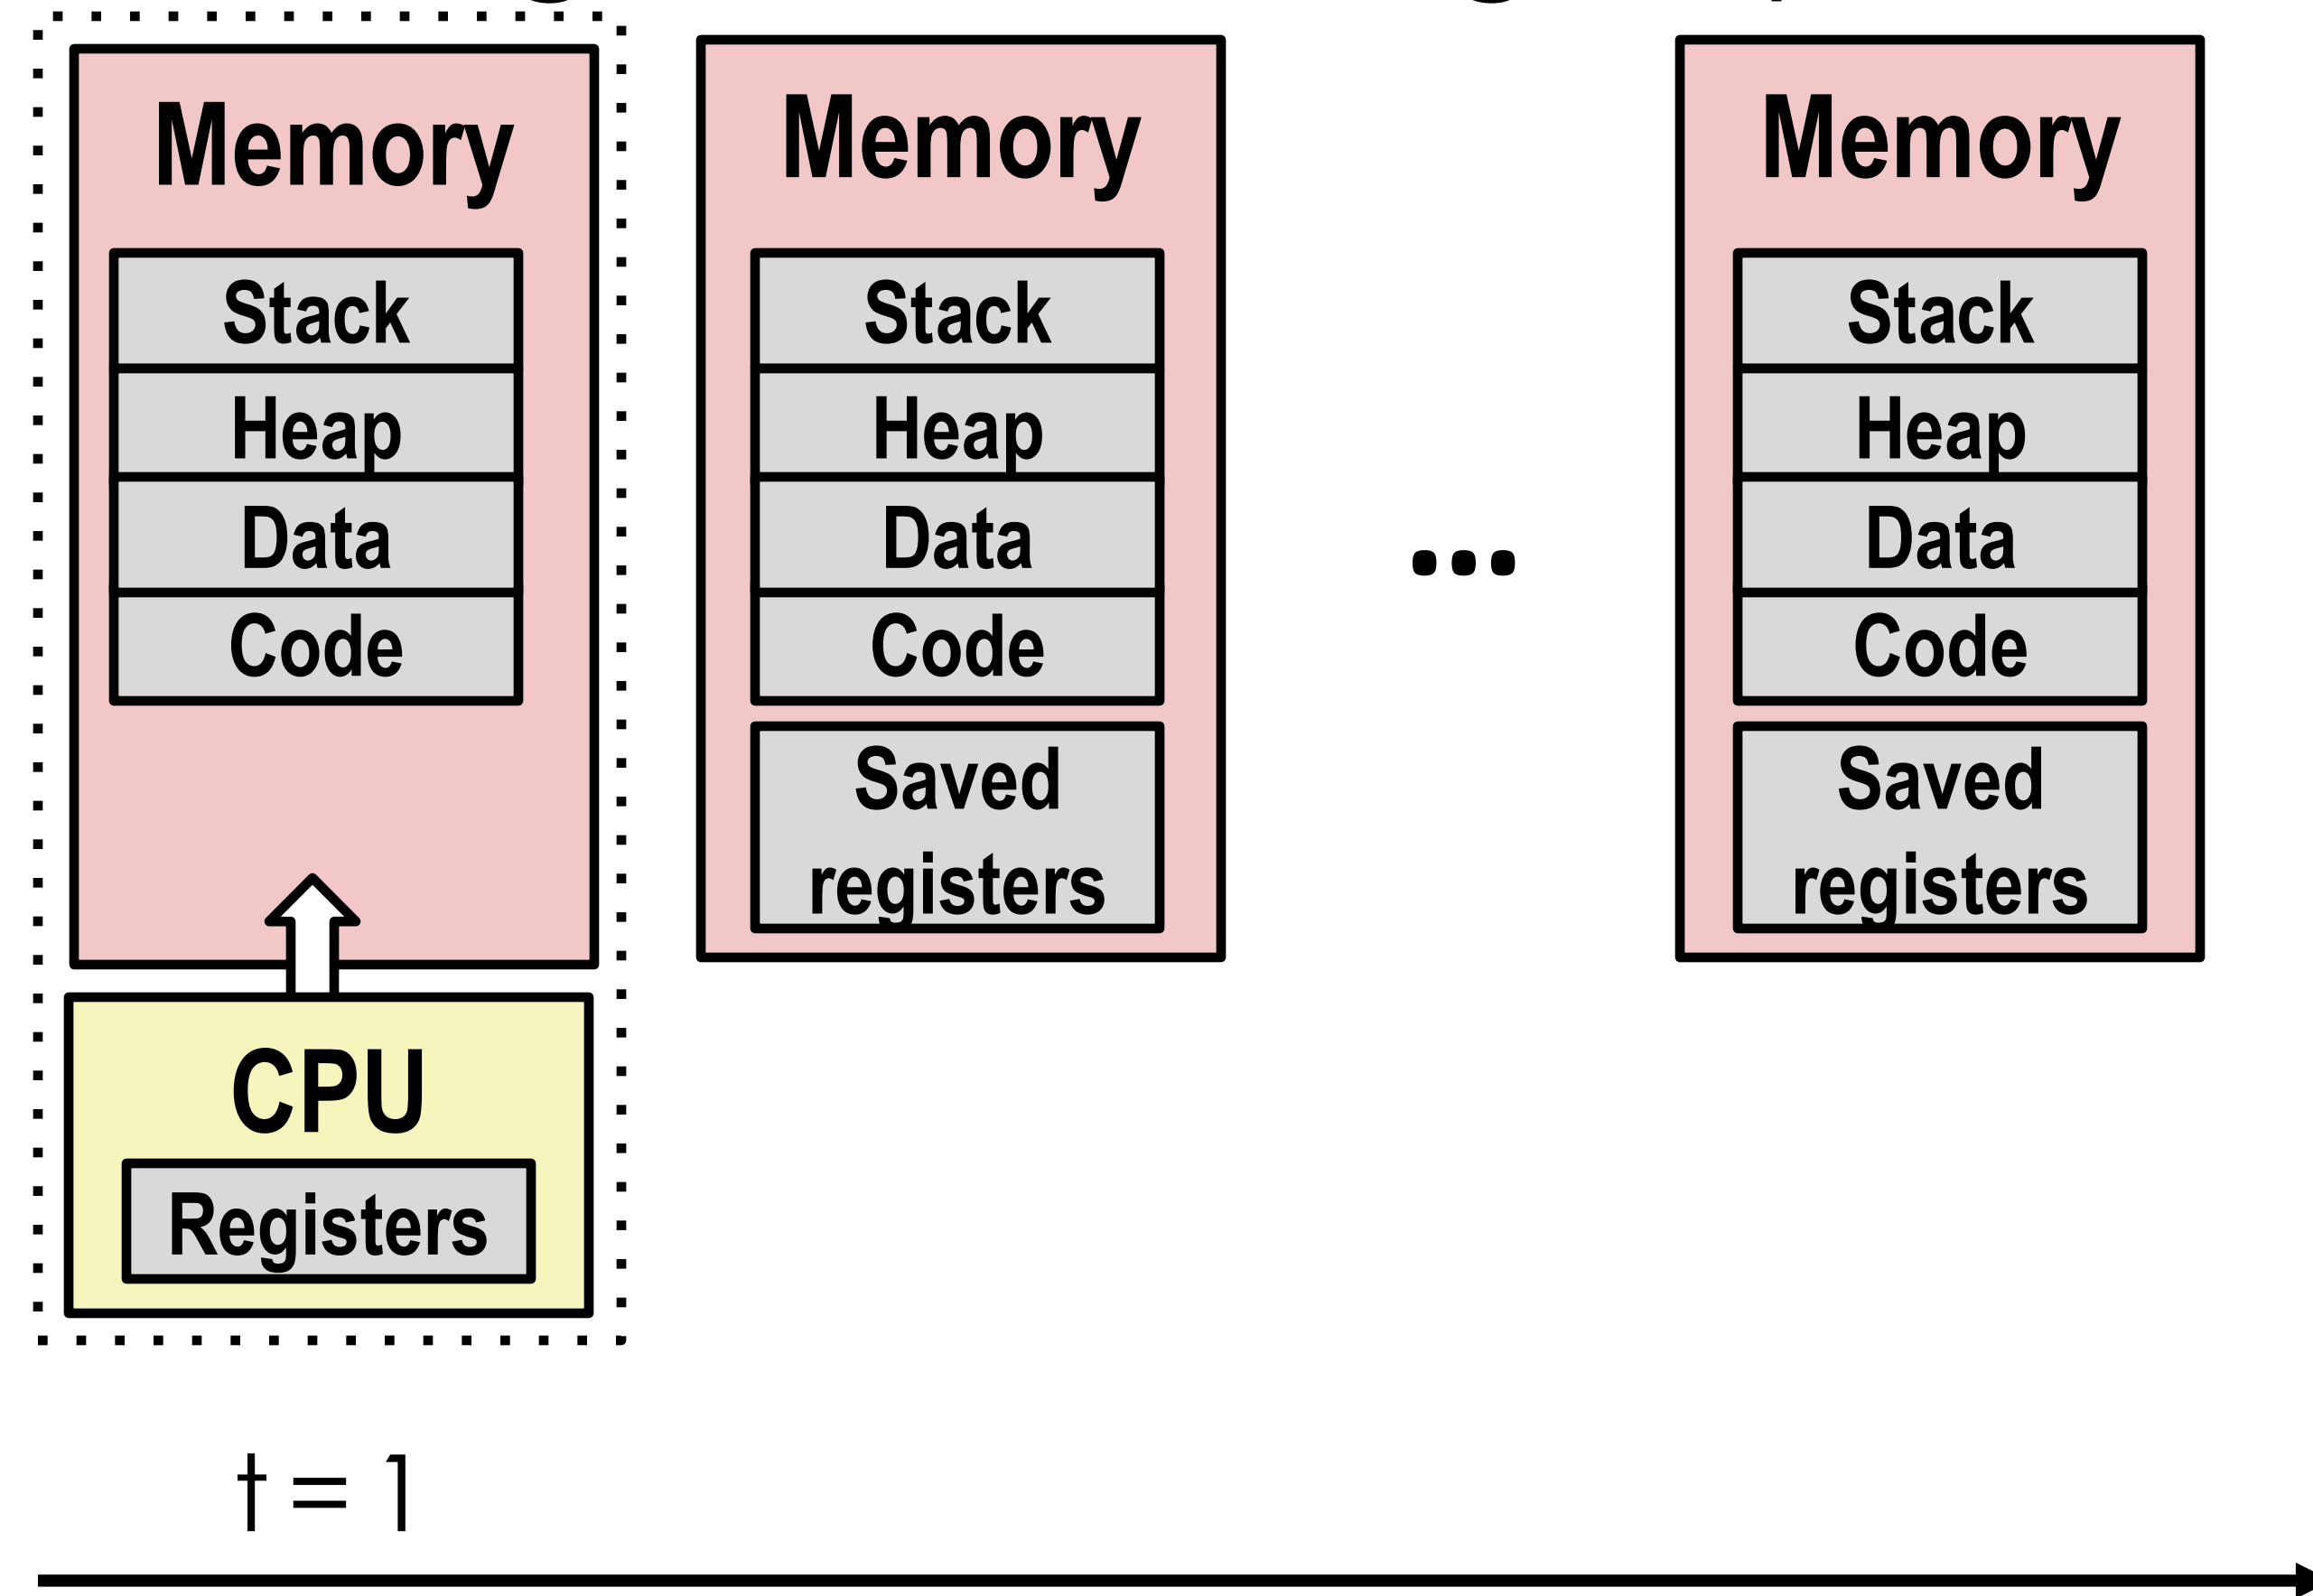
!!!we are N times slower when running N processes

# Multiprocessing: Time sharing of processors



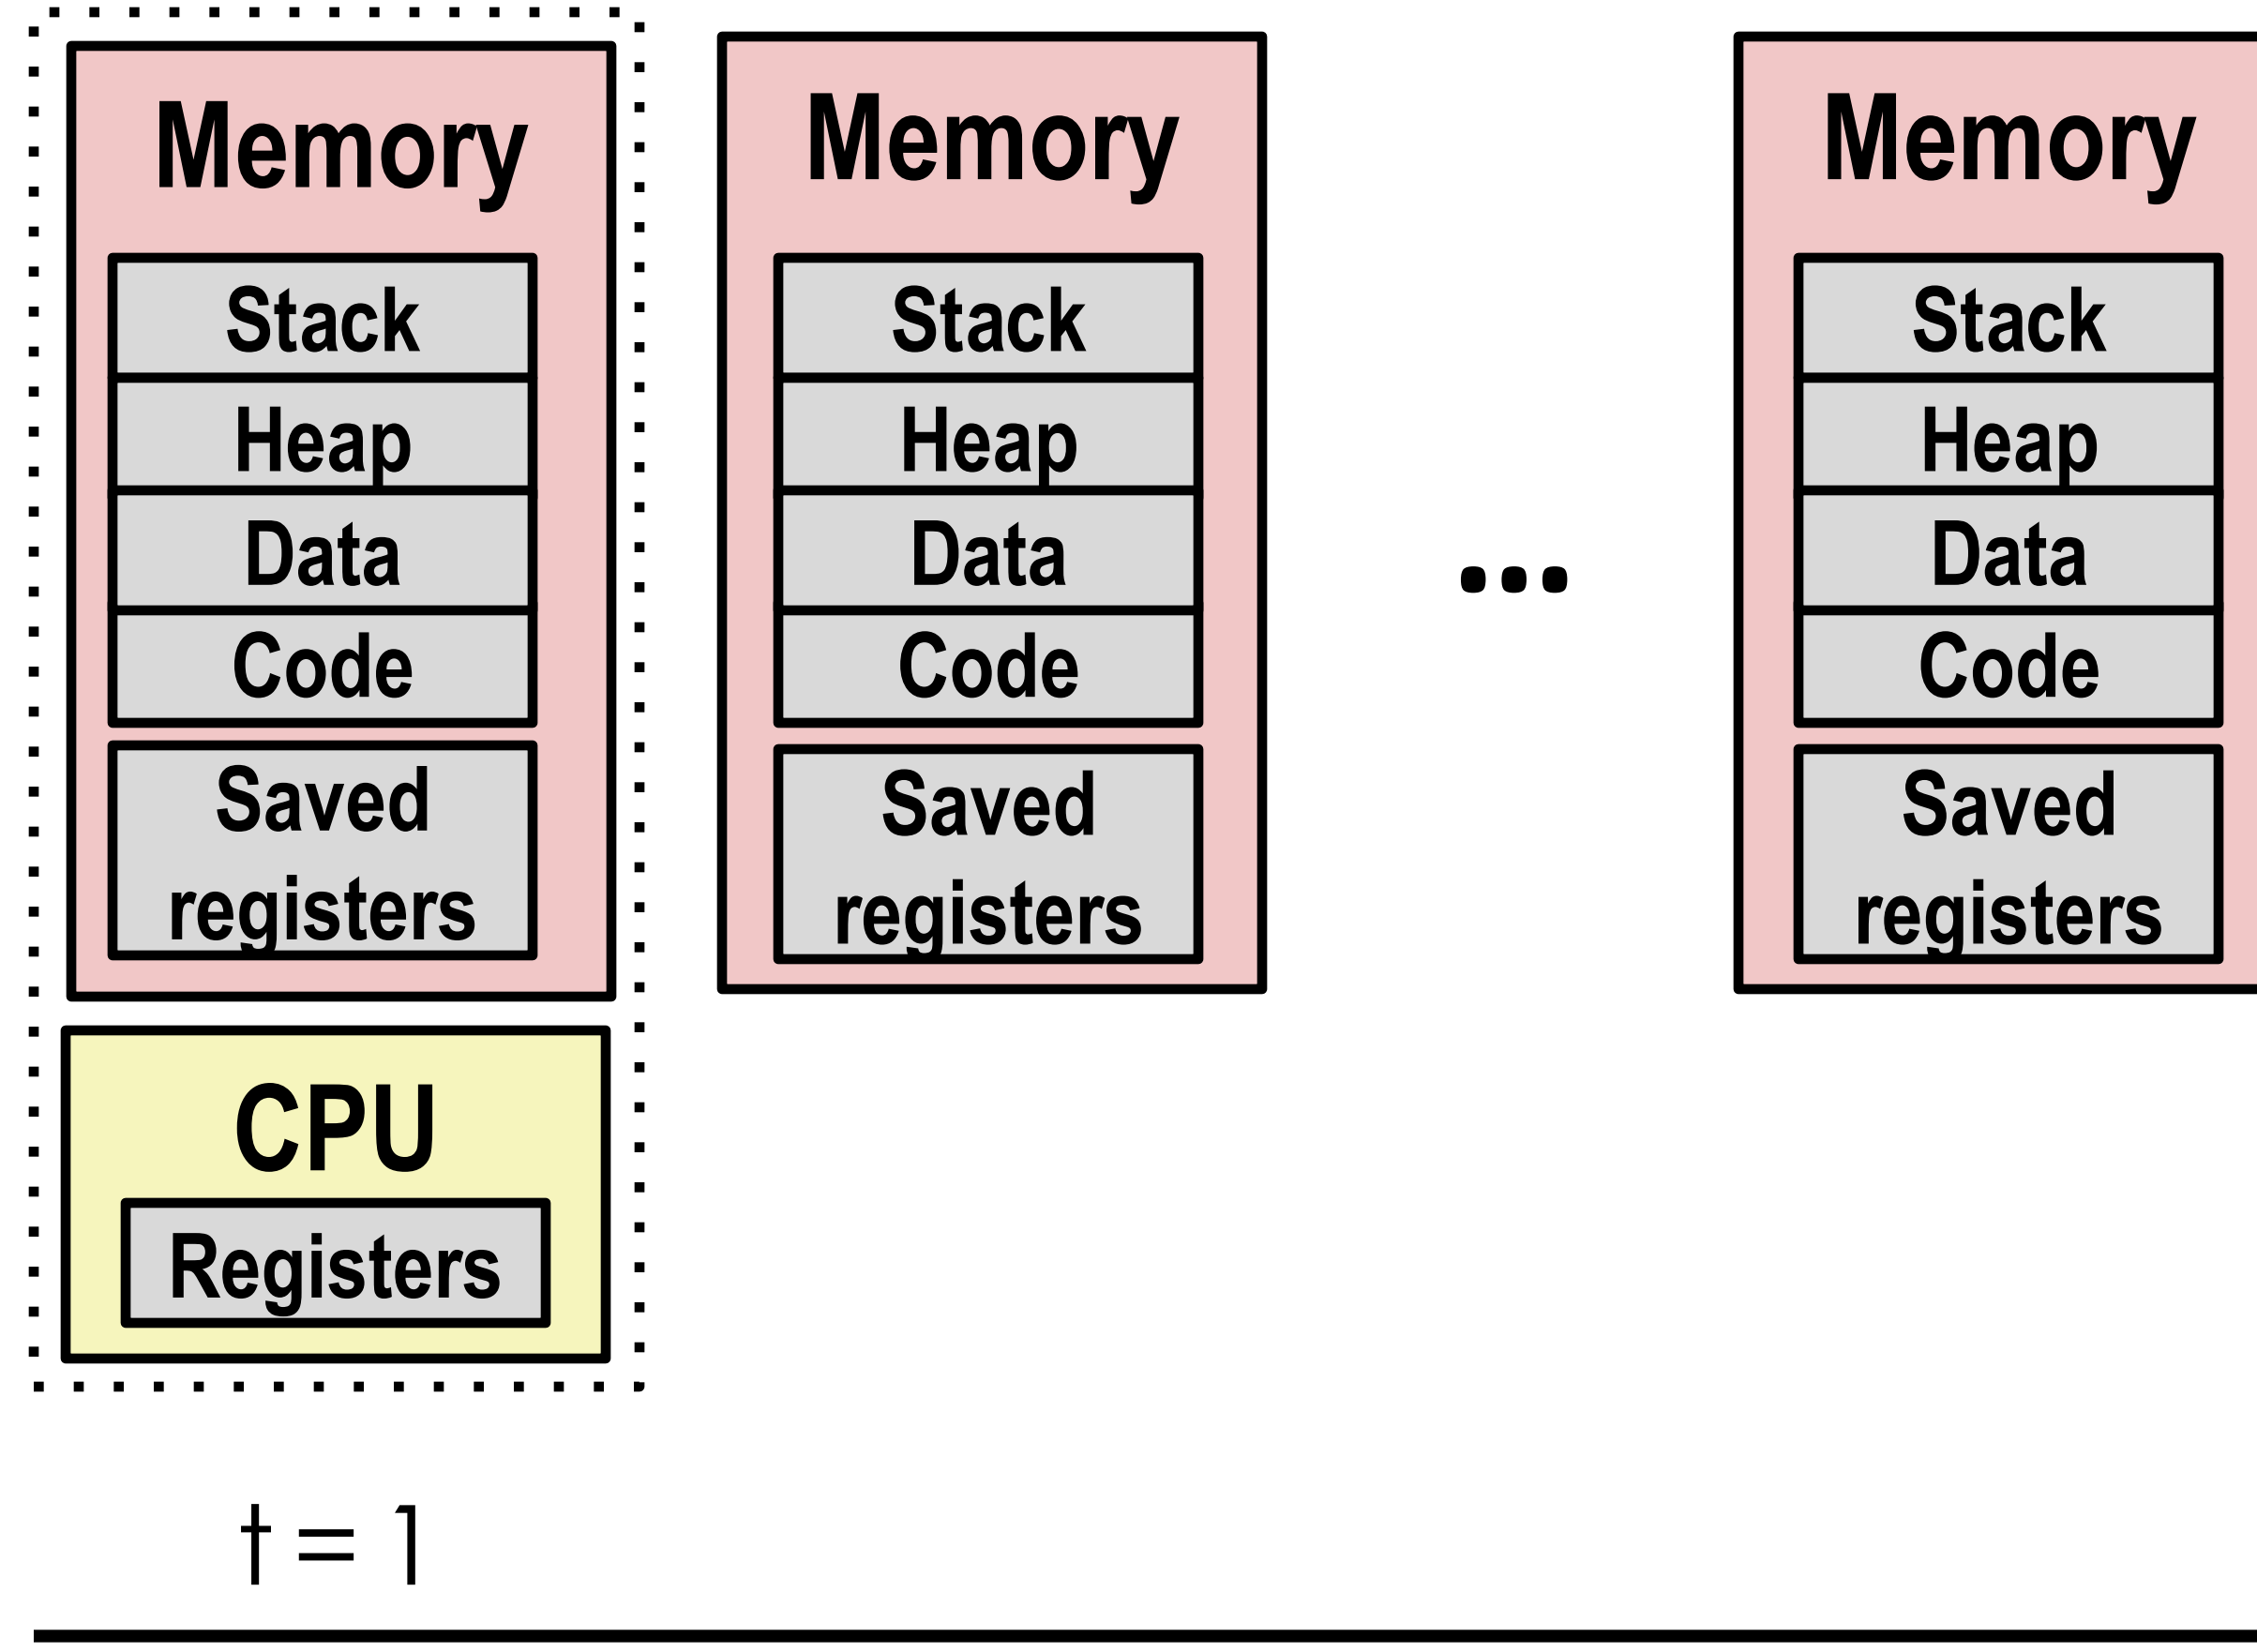
- Idea: Virtualize the CPU time as time slices
- Assign time slices to different processes

# Multiprocessing: Time sharing of processors



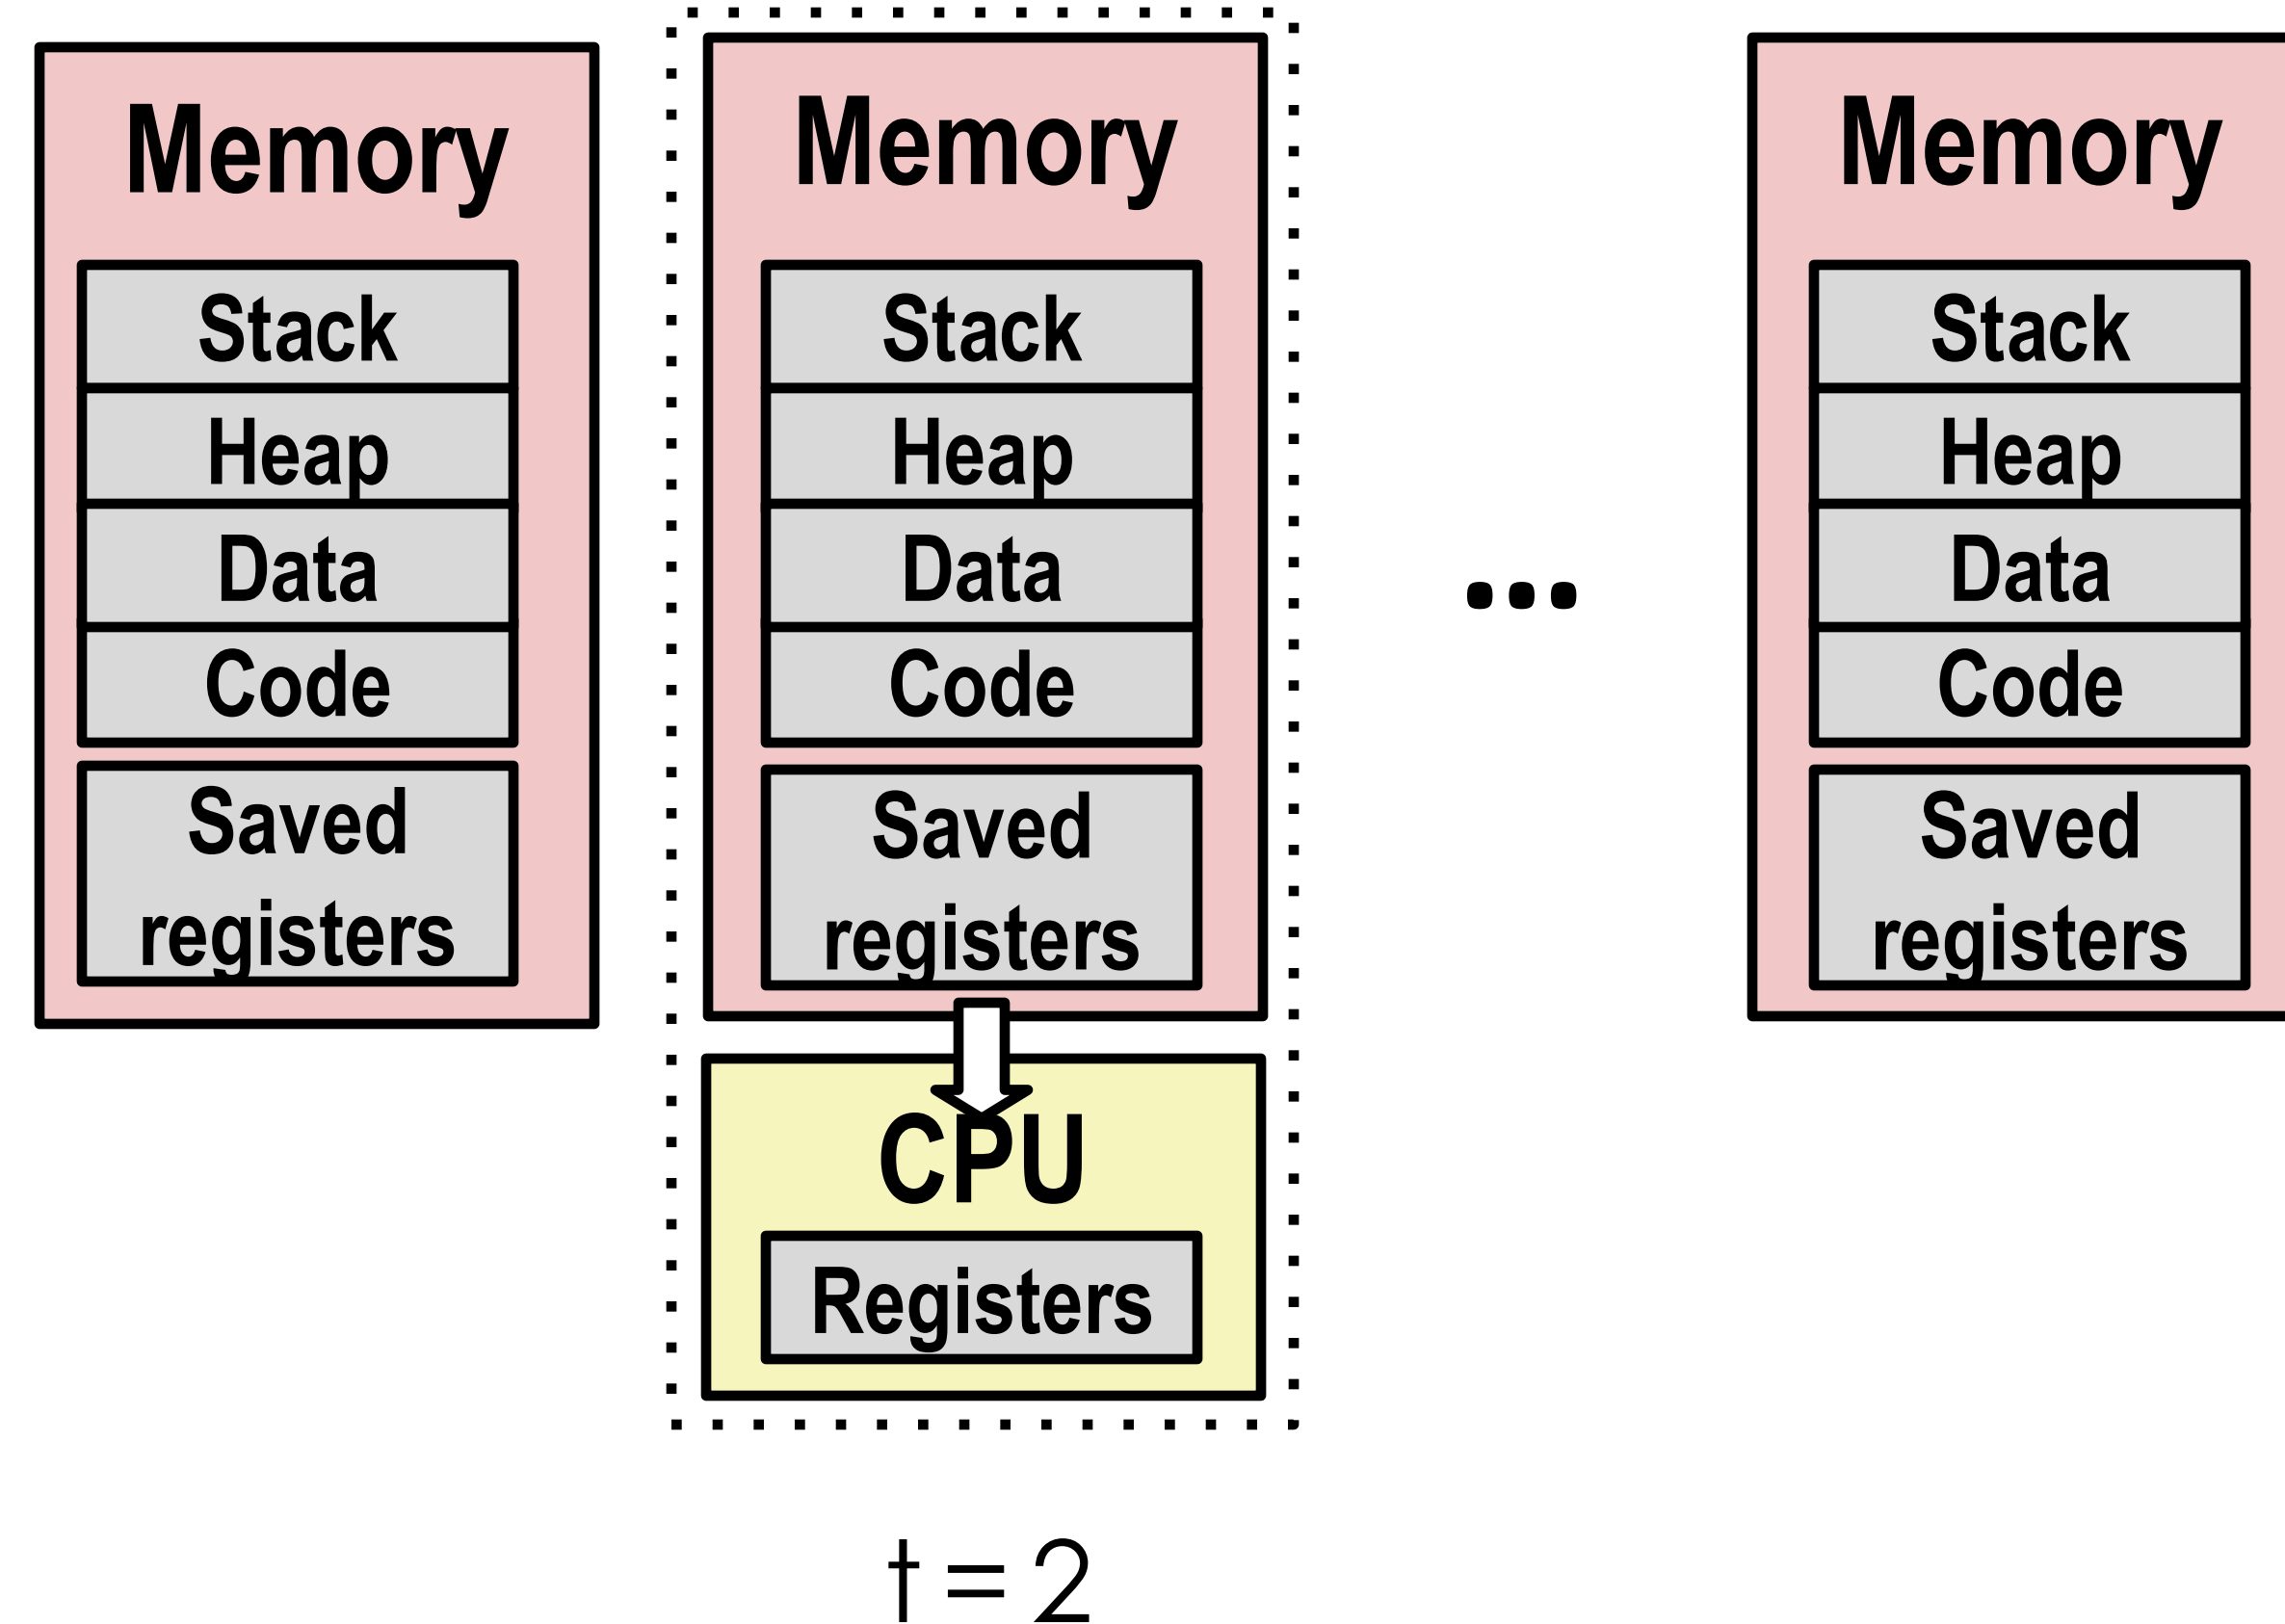
- Save current registers in memory

# Multiprocessing: Time sharing of processors



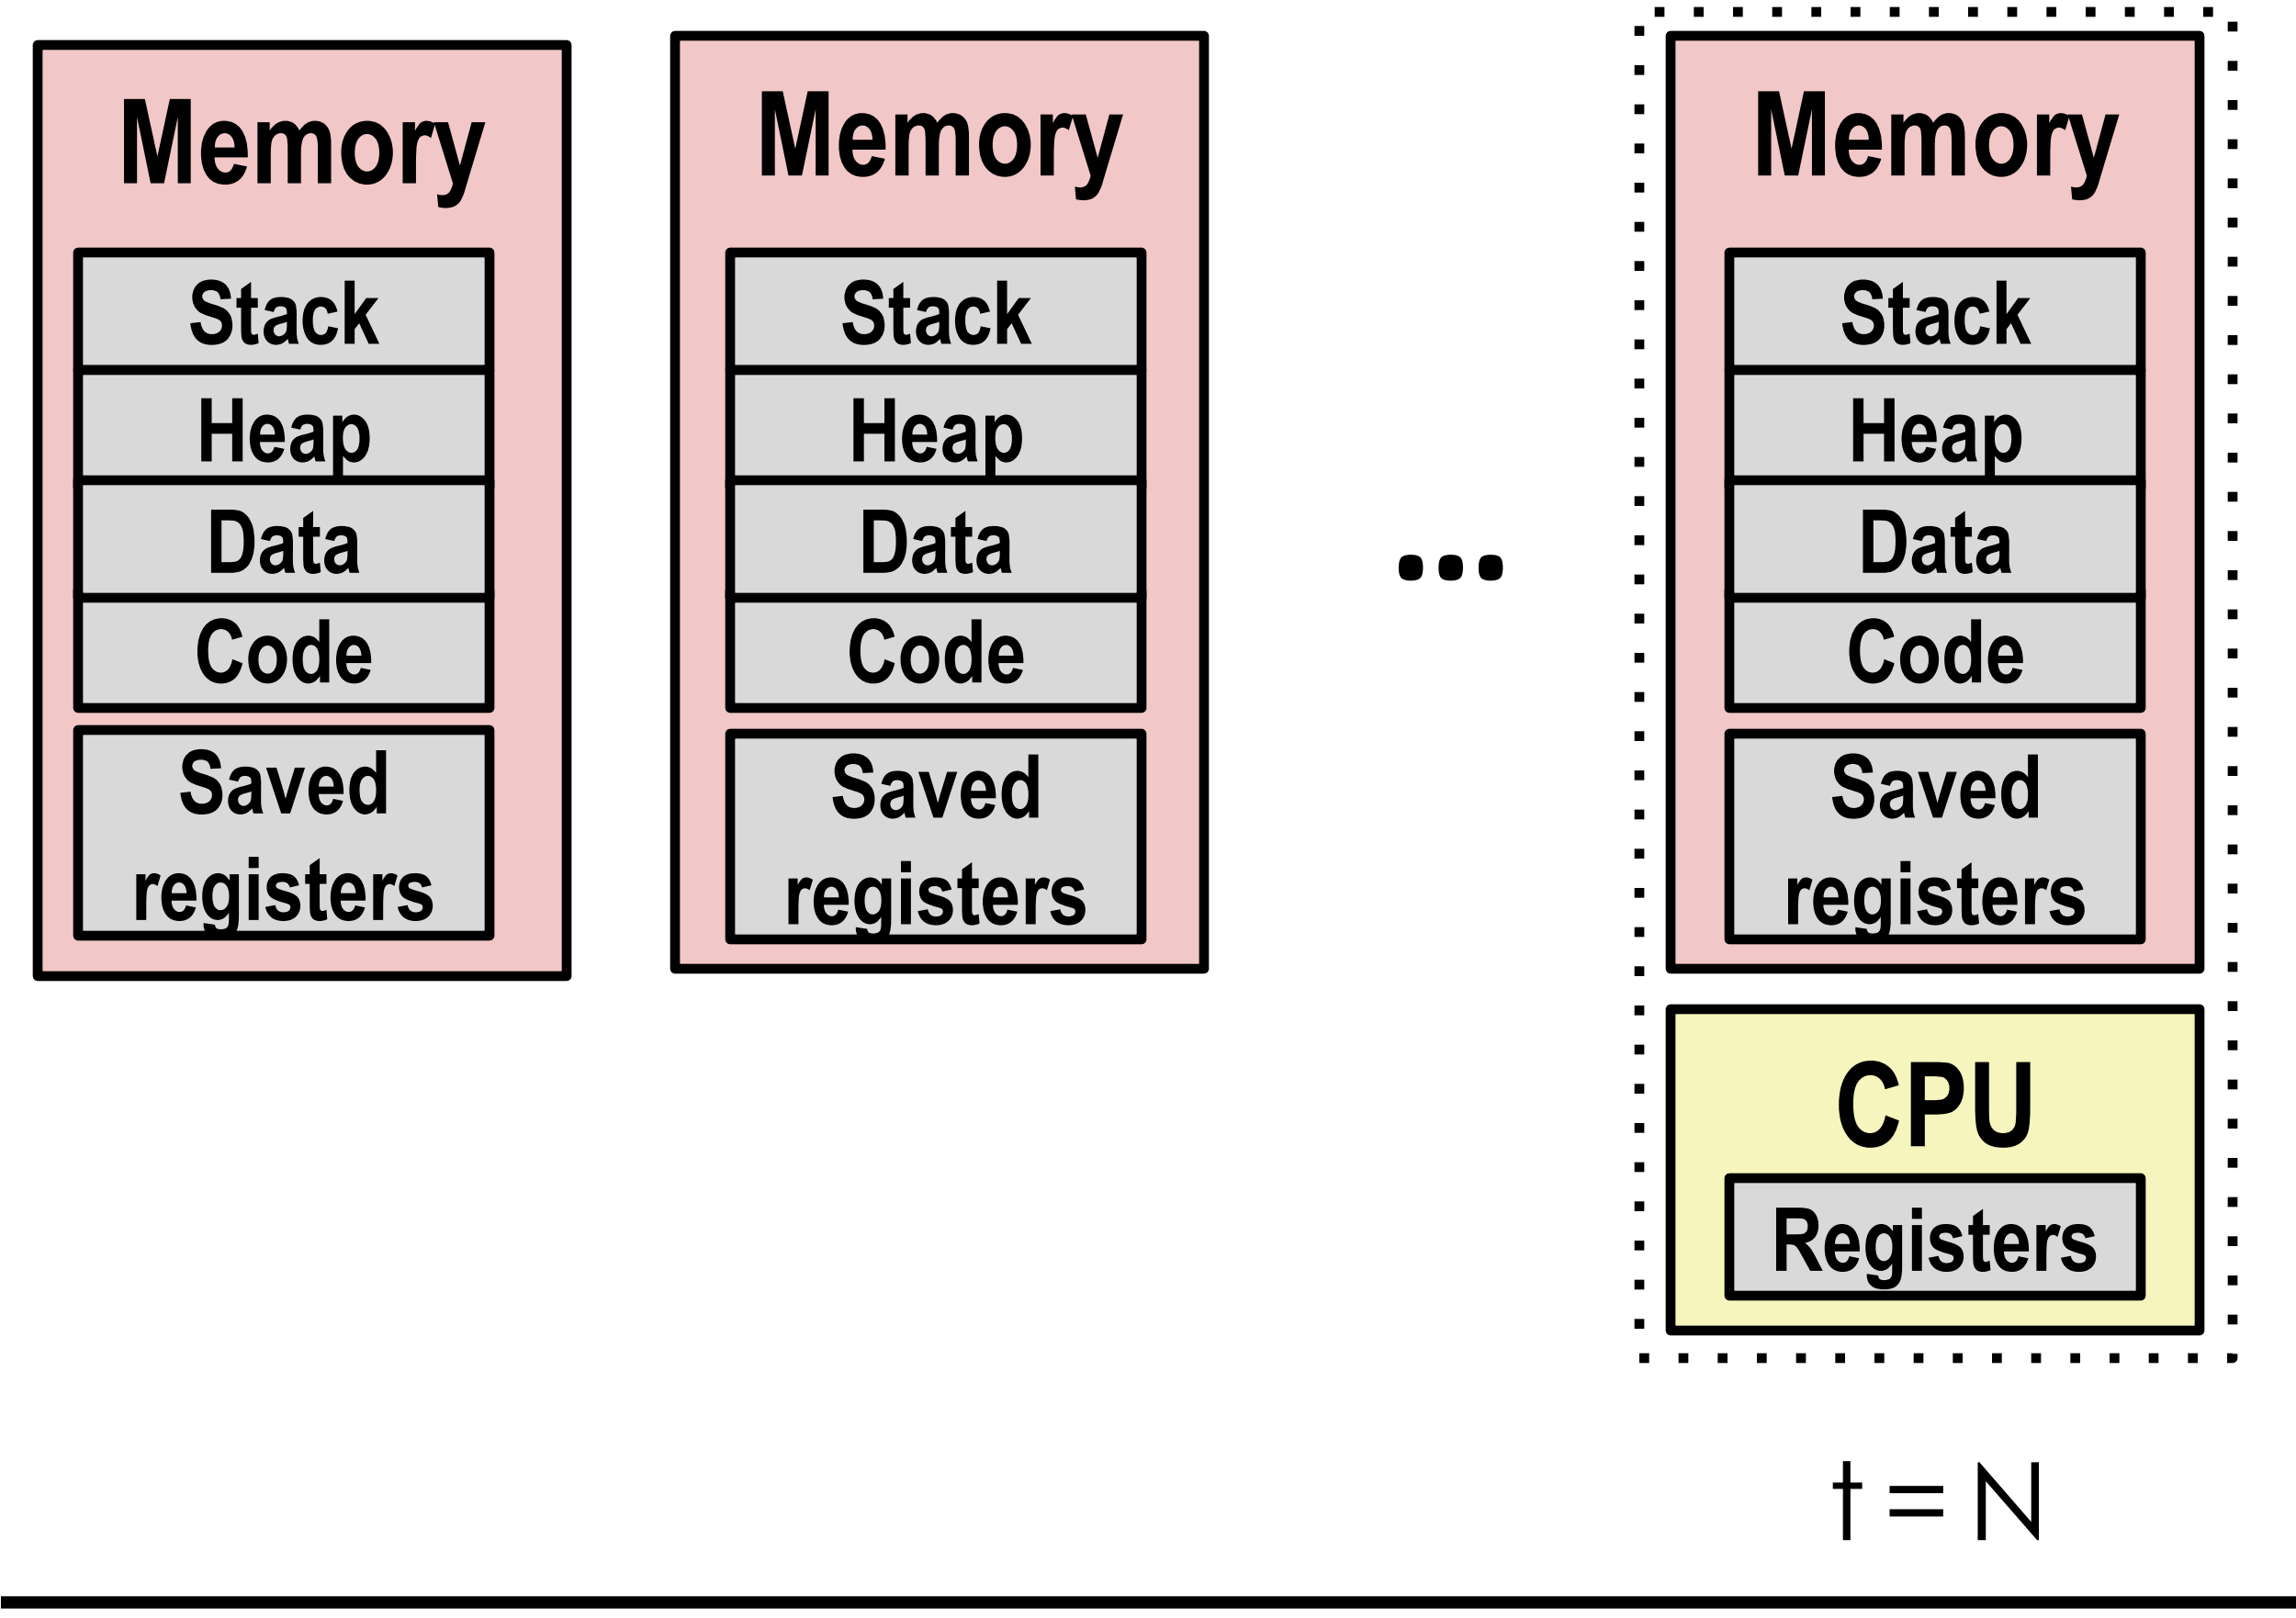
- Save current registers in memory

# Multiprocessing: Time sharing of processors



- Assign time slice  $t = 2$  to the next process
- Resume progress: Move Saved registers from memory to CPU

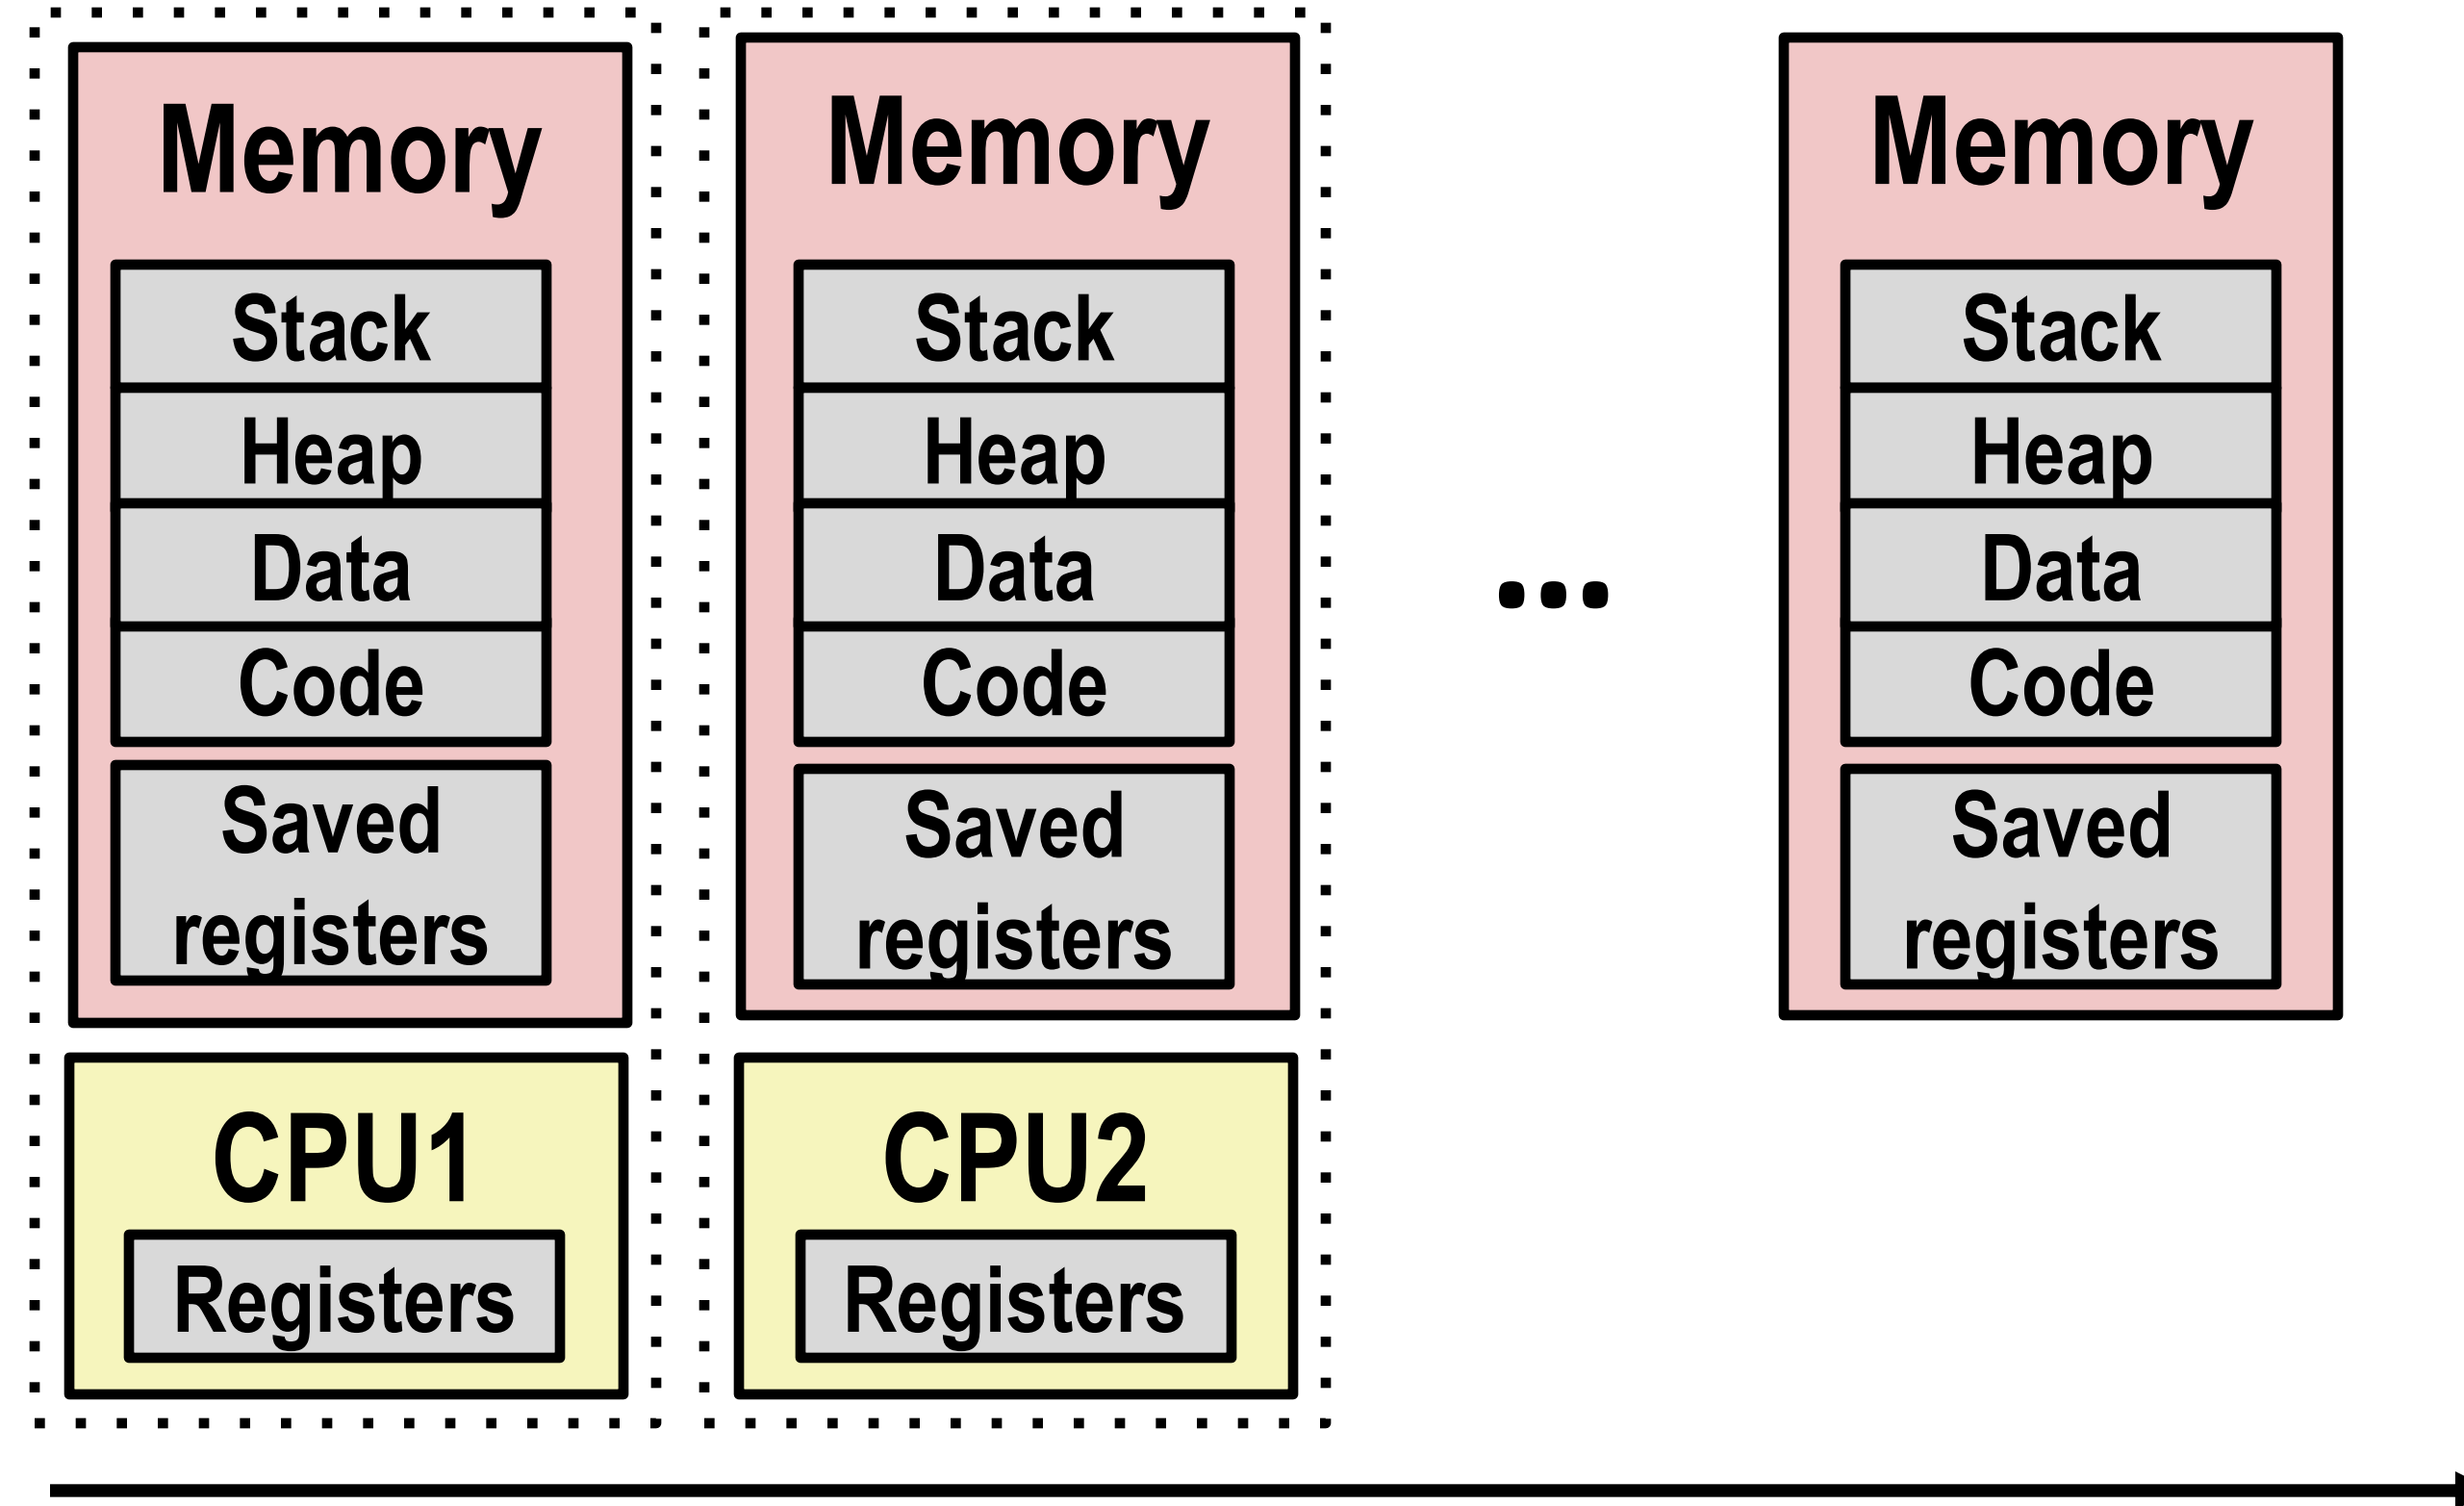
# Multiprocessing: Time sharing of processors



- Then we repeat.
- This is called **context switch**



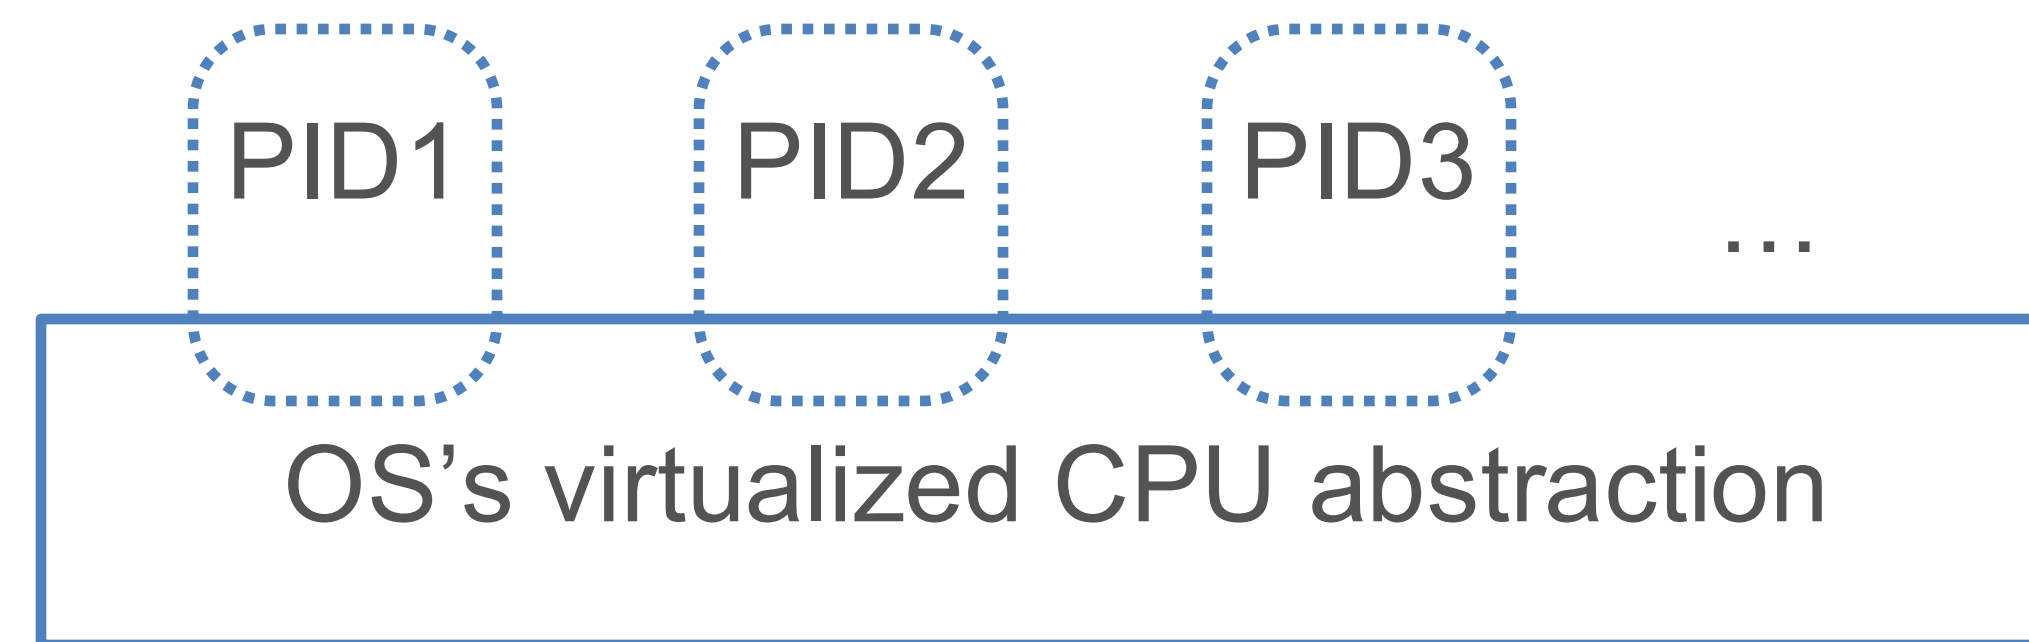
# Multiprocessing: Time sharing of multiple processors



Multiple CPU cores?

1. All processors sweep from left (1<sup>st</sup> process) to right (last process)
2. Each process accounts for  $\frac{1}{2}$  of the processes

# Let's Implement It!



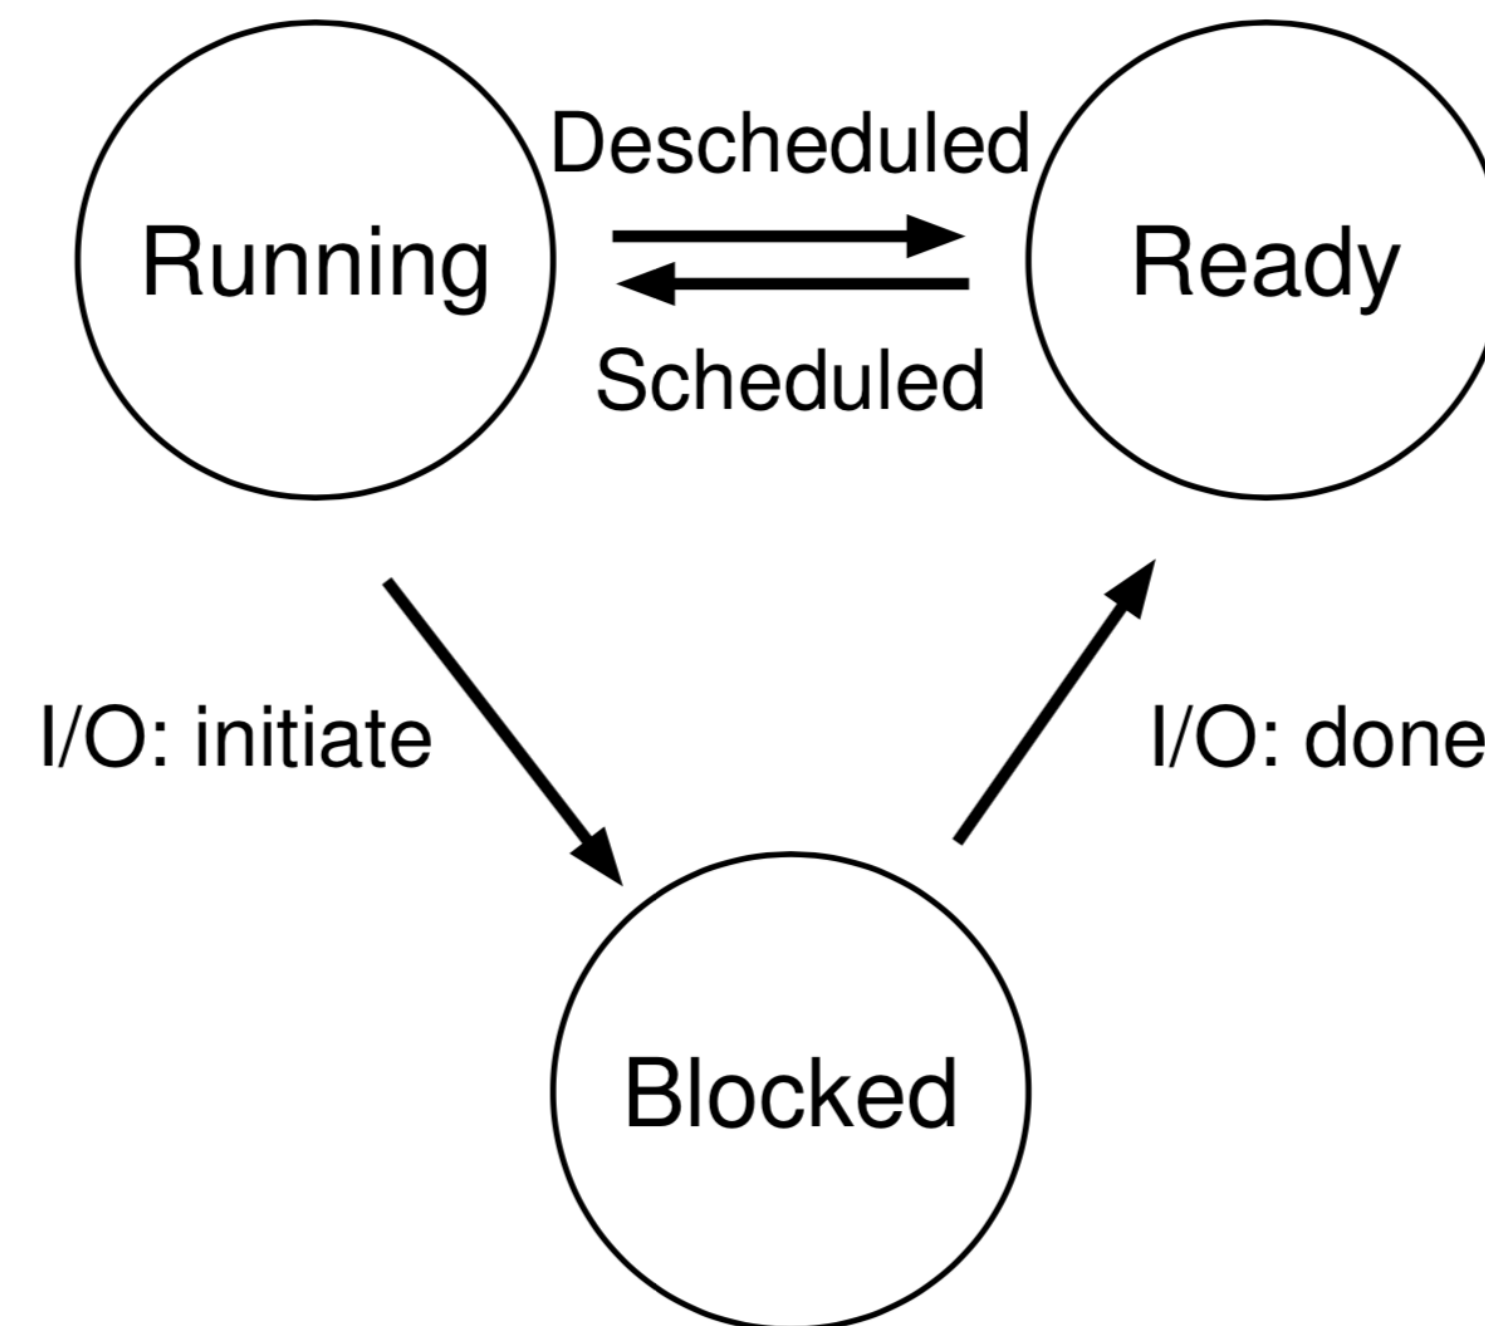
GAP1: How to virtualize CPU resources **temporally** and **spatially**?



Physical  
Processor

# Temporal Abstraction: Process State and CPU Time

- ❖ OS keeps moving processes between 3 states:



- ❖ Gantt Chart: A viz. to show what process runs when (on processor)



Scheduling question naturally emerges:

Q: how to schedule processes on time axis so **the objective** is optimal?


# Scheduling Policies/Algorithms

- Schedule: Record of what process runs on each CPU when
- Policy controls how OS time-shares CPUs among processes
- Key terms for a process (aka job):
  - **Arrival Time**: Time when process gets created
  - **Job Length**: Duration of time needed for process
  - **Start Time**: Time when process first starts on processor
  - **Completion Time**: Time when process finishes/killed
  - **Response Time** = Start Time — Arrival Time
  - **Turnaround Time** = Completion Time — Arrival Time
- Workload: Set of processes, arrival times, and job lengths that OS Scheduler has to handle

# Scheduling Policy: FIFO

- ❖ First-In-First-Out aka First-Come-First-Serve (FCFS)
- ❖ Ranking criterion: Arrival Time; no preemption allowed

**Example:** P1, P2, P3 of lengths 10,40,10 units arrive closely in that order

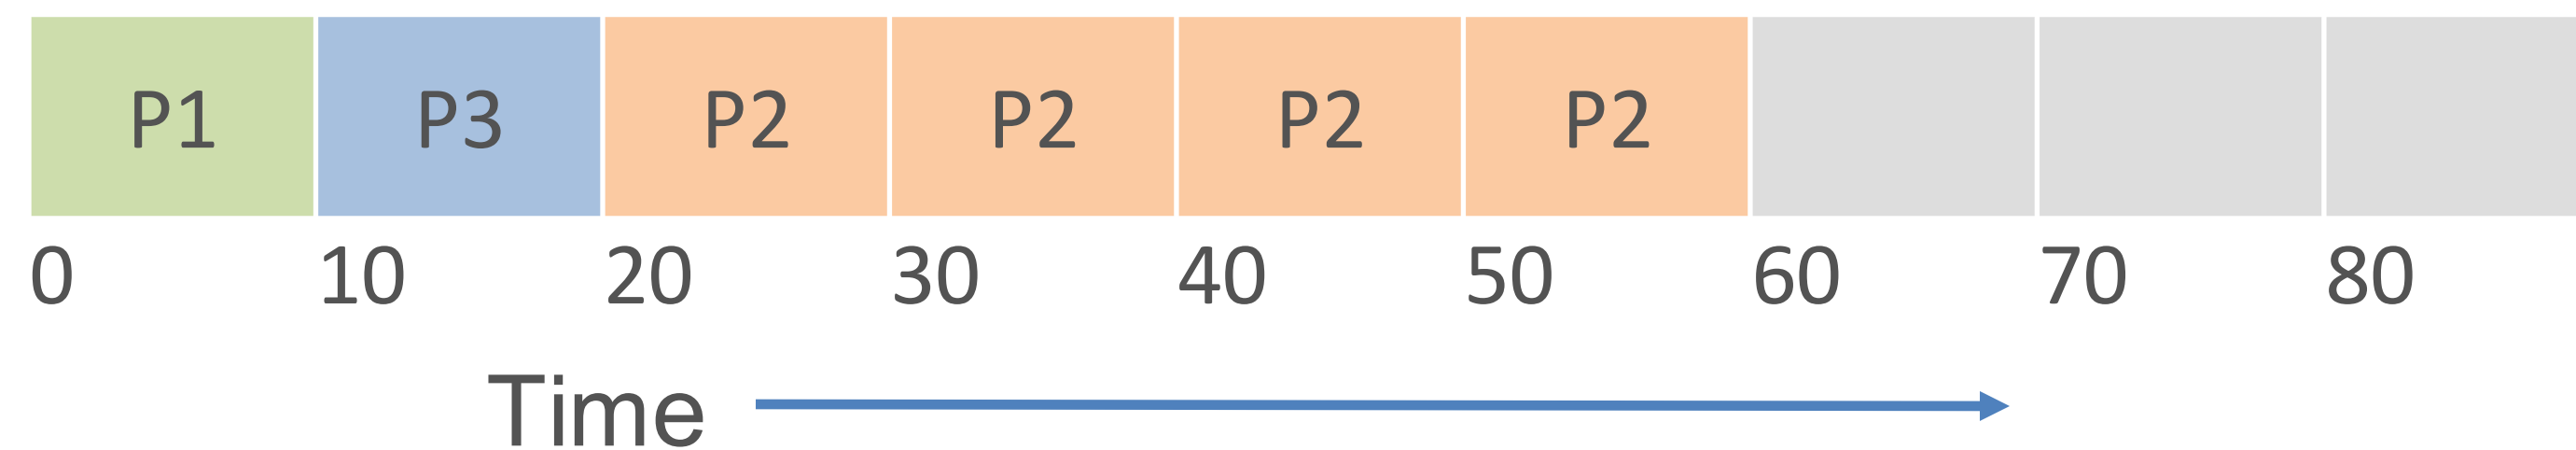


- ❖ Main con: Short jobs may wait a lot, aka “Convoy Effect”

# Scheduling Policy: SJF

- ❖ Shortest Job (next) First
- ❖ Ranking criterion: Job Length; no preemption allowed

**Example:** P1, P2, P3 of lengths 10,40,10 units arrive closely in that order




Process	Arrival Time	Start Time	Completion Time	Response Time	Turnaround Time
P1	0	0	10	0	10
P2	0	20	60	20	60
P3	0	10	20	10	20
(FIFO Avg: 20 and 40)				Avg: 10	30

- ❖ Main con: Not all Job Lengths might be known beforehand
- ❖ Long processes may be held off indefinitely

# Example Exam Q1: Round Robin Schedule

- ❖ RR does not need to know job lengths
- ❖ Fixed time *quantum* given to each job; cycle through jobs

**Example:** P1, P2, P3 of lengths 10,40,10 units arrive closely in that order




- ❖ RR is often very fair, but Avg Turnaround Time goes up!

# Example Exam Q2: SCTF

- ❖ Shortest Completion Time First
- ❖ Jobs might not all arrive at same time; preemption possible

**Example:** P1, P2, P3 of lengths 10,40,10 units arrive at different times






# Scheduling Policies/Algorithms

- In general, not all Arrival Times and Job Lengths will be known beforehand. But preemption is possible.
- **Key Principle: Inherent tension in scheduling between overall workload performance and allocation fairness**
  - Performance metric is usually *Average Turnaround Time*
  - Many fairness metrics exist, e.g., Jain's fairness index
- 100s of scheduling policies studied! Well-known ones: FIFO, SJF, STCF, Round Robin, Random, etc.
  - Different criteria for ranking; preemptive vs not
  - Complex “multi-level feedback queue” schedulers
  - ML-based schedulers are “hot” nowadays!

# Scheduling in ChatGPT



- What is the response time
- What is the turnover time
- What is fairness?
- Do we know the job length?
- Can we run S1/S2/S3 together?
- How to schedule?